

HealthLedger AI

Patent, Trademark, Trade Dress & Copyright Disclosure

AI-powered billing fraud detection and accounting compliance for hospital finance departments. Grounded in doctoral research. Built through on-site immersion. Custom-trained on hospital cooperation agreements.

Inventor:	Maria Polychroniadou
Contact:	polychroniadou@gmx.de
Brand:	HealthLedger AI
Domain:	healthledgerai.com
Contact Email:	info@healthledgerai.com
Date of Conception:	Saturday 14 March 2026
First Reduction to Practice:	Tuesday 17 March 2026 (Version 1.0 — working system)
Phase 2 Engineering:	28 March 2026 (fault tolerance architecture designed)
Disclosure Date:	June 26, 2026
Document Version:	4.0 — FINAL · Complete Code + Fault Architecture Disclosure
Status:	CONFIDENTIAL — Pre-Filing — Attorney Review

ATTORNEY-CLIENT PRIVILEGED AND WORK PRODUCT PROTECTED. Contains source code, fault-tolerance architecture, trade secrets, and strategic IP. Prepared exclusively for patent counsel. Unauthorised reproduction prohibited.

TABLE OF CONTENTS

- A. Brand Identity, Trademark & IP Overview
 - 1. Field of the Invention
 - 2. Background, Motivation & Academic Basis
 - 3. Conception Timeline & Reduction to Practice
 - 4. Summary of the Invention — 12-Feature Detection System
 - 5. Detailed Technical Description
 - 5.1 System Architecture (5-Layer Model)
 - 5.2 12-Feature Billing Fraud Detection Engine
 - 5.3 AI Agent & Session Architecture
 - 5.4 Cooperation Agreement RAG Pipeline
 - 5.5 Domain Training & On-Site Immersion Module
 - 5.6 Fault Tolerance & Remote Access Architecture (Phase 2)
 - 5.7 Frontend Interface & Deployment
 - 5.8 Visual Design System
 - 6. Source Code Listings
 - 6.1 Core AI Agent Module (managed-agents-starter.ts)
 - 6.2 12-Feature Detection Engine (lib/detection-engine.ts)
 - 6.3 Fault Tolerance Architecture (lib/fault-tolerance.ts)
 - 6.4 RAG Pipeline (lib/ingestion-pipeline.ts)
 - 6.5 Domain Training Module (lib/domain-training.ts)
 - 6.6 Frontend + Deployment Config (app/page.tsx + next.config.js + CNAME)
 - 7. Prior Art Analysis
 - 8. Patent Claims (16 Draft Claims)
 - 9. Academic References
 - 10. Abstract
 - 11. Inventor Declaration

A. BRAND IDENTITY, TRADEMARK & IP OVERVIEW

IP Type	Subject Matter	Action Required
Patent	12-feature detection system, AI compliance engine, RAC pipeline, → PCT/national phase fault tolerance architecture, on-site methodology, UI in (EURO, USPTO, DPMA)	
Trademark	'HealthLedger AI', 'healthledgerai.com', 'Healthcare Accounting Intelligence', taglines, logo dot (NICE Class 42)	EUIPO + USPTO + DPMA
Trade Dress	Colour palette, lily botanical, diamond motif, pulsing dot, typography pairing	Include in TM application as design mark
Copyright	All source code (6 listed files), SVG botanical artwork, CSS design system	Register with copyright office for statutory damages
Trade Secret	On-site immersion methodology details, institutional training data, custom rules engine	NDA with all hospital clients before discovery sessions

Trademark Marks Claimed

Primary Mark:	HealthLedger AI
NICE Class:	42 — AI software; compliance software; SaaS for healthcare finance
Domain:	healthledgerai.com (CNAME record confirmed in repository)
Email:	info@healthledgerai.com
Secondary Mark 1:	"Healthcare Accounting Intelligence" — navigation subtitle
Secondary Mark 2:	"Something important is being built." — footer tagline
Secondary Mark 3:	"The intelligence healthcare accounting has been waiting for."
Logo Element:	Pulsing sage dot (animated, 2.8s cycle) adjacent to logotype
First Use Date:	14 March 2026 (date of conception; repository timestamps confirm)

Trade Dress — Visual Identity

	Cream #F4EFE4	Primary background — warm parchment
	Gold #A08558	Primary accent — antique gold, all highlights and rules
	Sage #4A6B5A	Secondary accent — logo dot, botanical, italic headline
	Ivory #EDE7D8	Secondary surface — contact section background
	Blush #D4A9A0	Tertiary — rare, particle field only
	Ink #1E1A14	Near-black warm — all primary text
■ Original SVG Lily Botanical: Botticelli-inspired lily (3 blooms, 3 leaves, curved stem). Sage fill (#4A6B5A), ivory/cream petals, gold stamens. Parallax-mounted at translateZ(-100px). Both copyright and trade dress element.		

- **Diamond Ornament Motif:** 45°-rotated square used throughout as brand signature — floating ornaments, section dividers, step accents. Rendered in gold (#C9AA7C).
- **Pulsing Logo Dot:** 8px sage circle with breathing box-shadow pulse (2.8s). Positioned left of logotype. Functions as the animated logo mark.
- **Typography:** Cormorant Garamond serif + Jost sans-serif pairing. Weights 300–500 only.
- **Gold Hairline Rules + Grain Texture + Radial Ambient Gradients.**

Copyright Coverage — All Source Files

Copyright:	© 2026 HealthLedger AI — Maria Polychroniadou. All rights reserved.
Source Files:	managed-agents-starter.ts · lib/detection-engine.ts · lib/fault-tolerance.ts · lib/ingestion-pipeline.ts · lib/domain-training.ts · app/page.tsx · app/globals.css · app/layout.tsx · next.config.js · components/Botanical.tsx · components/IntroOverlay.tsx · components/ParticleField.tsx · components/HeroScene.tsx · components/FloatingOrbs.tsx · components/CustomCursor.tsx · components/Marquee.tsx · components/ProblemHook.tsx · components/SectionHow.tsx · components/SectionContact.tsx
Artwork:	SVG Lily Botanical (Botanical.tsx) — original vector artwork
Evidence:	Git repository (mrply97/AI, GitHub) — timestamped commit history

1. FIELD OF THE INVENTION

The present invention relates to artificial intelligence systems for healthcare financial integrity and accounting compliance. More particularly, the invention concerns: (i) a 12-feature billing fraud and compliance detection engine for hospital accounting departments; (ii) a managed AI agent system trained on hospital cooperation agreement corpora providing real-time clause-traceable compliance guidance; (iii) a fault-tolerant remote-access architecture for safe deployment on live hospital information systems; and (iv) a novel on-site immersion methodology for capturing institutional accounting knowledge as AI training data.

2. BACKGROUND, MOTIVATION & ACADEMIC BASIS

Academic Foundation

The invention is grounded in doctoral research at European University Cyprus investigating how AI can transform accounting processes in hospitals. The PhD research documents billing error patterns, fraud typologies, and compliance gaps across hospital finance departments. The research provided both the evidence base for the product specification and the methodology for hospital engagement. Studies consulted establish that up to 80% of medical bills contain at least one error (Medical Billing Advocates of America, 2024), with manual billing error rates ranging from 7% to 26% (Nimkar et al., 2024). Healthcare billing fraud costs an estimated USD 68 billion annually in the US alone (IJSRA, 2024).

Motivating Real-World Incident

The direct inspiration for the patient identity collision detection feature (Feature 3) was a real incident reported by a colleague at the European Interbalkan Medical Center, Thessaloniki, Greece: two patients shared the same first name, last name, and date of birth. An invoice was assigned to the wrong patient, and that patient faced legal proceedings as a result of the data entry error. HealthLedger AI's identity collision detection with parental name verification (the Greek secondary identifier) was designed specifically in response to this incident.

Technical Problem Statement

- **Volume of Cooperation Agreements.** A hospital network maintains 200+ active cooperation agreements, each containing accounting obligations (revenue recognition, cost allocation, reconciliation, penalties). Total actionable obligations: thousands per network.
- **Billing Fraud Typologies.** Six principal fraud types documented in the literature: upcoding, phantom billing, unbundling, duplicate invoicing, surprise billing, and ghost patient billing. No prior art system detects all six in a single automated pipeline.
- **Inconsistent Clause Interpretation.** Agreement clauses are interpreted differently by different accounting staff. A single misread clause generates €10k–€500k+ liability per period.
- **No Domain-Specific AI Exists.** General-purpose AI assistants are not trained on hospital cooperation agreement corpora and cannot produce clause-cited, auditable accounting compliance determinations.

■ **Patient Data Integrity Risk.** Any AI system connecting to live hospital information systems must guarantee that it never modifies source data and can recover gracefully from failure without corrupting patient or billing records. No prior art system implements the specific fault-tolerance architecture designed for this purpose.

3. CONCEPTION TIMELINE & REDUCTION TO PRACTICE

ATTORNEY NOTE — PRIORITY DATE: The conception date is 14 March 2026. The first working reduction to practice is 17 March 2026. File the provisional application BEFORE the website healthledgerai.com becomes publicly accessible. European Patent Convention (EPC) provides NO grace period — any public disclosure before filing is an absolute bar. Under US AIA a 12-month grace period applies from inventor's own disclosure. Recommend simultaneous provisional filing in US and EP.

Date	Event	Evidence
14 March 2026	Date of Conception First Python program written (Saturday) Zero prior coding knowledge; first code on Replit	Developer Journal (personal record) Replit session logs
14–17 March 2026	Active Development — Version 1.0 12 detection features built and tested Python → production-grade TypeScript upgrade	Developer Journal — detailed lesson-by-lesson account Git commit history
17 March 2026	First Reduction to Practice (Tuesday) Working system: reads Excel, detects 11 fraud types writes HealthLedger_Alerts.xlsx — complete pipeline	Developer Journal: 'Version 1.0 is complete' Git repository
28 March 2026	Phase 2 Engineering Requirements Documented Fault tolerance architecture designed (read-only, checkpointing, integrity hashing, remote access)	Developer Journal — Phase 2 section Podcast session record
March–June 2026	Upgrade to Managed AI Agents Architecture Next.js 14 frontend, Anthropic claude-opus-4-8K healthledgerai.com domain registered (CNAME)	Git repository (mrply97/AI) CNAME file in repository
26 June 2026	This Patent Disclosure Document Prepared	This document

4. SUMMARY OF THE INVENTION — 12-FEATURE DETECTION SYSTEM

HealthLedger AI is a multi-component AI system comprising: (A) a 12-feature billing fraud and compliance detection engine; (B) a cooperation agreement RAG compliance layer; (C) a fault-tolerant deployment architecture; and (D) a novel on-site immersion training methodology.

A — 12-Feature Detection Engine

The following table lists all 12 detection features, their fraud/error targets, and real-world financial impact. First implemented 14–17 March 2026.

#	Feature	What It Detects	Impact / Legal Basis
1	Duplicate Invoice Detection	Same invoice submitted twice (vendor + amount + date fingerprint)	Most common billing error; hospitals pay twice without detection
2	Risk Scoring (CRITICAL/HIGH/MEDIUM/LOW)	High/Medium/Low classification for alert triage prioritisation	Focus accounting staff on highest-value alerts first
3	Patient Identity Collision	Two patients: same first name + last name + DOB Invoice assigned to wrong patient	DOBs legal liability (real incident: European Interbalkan Medical Center, Thessaloniki)
4	CFO Dashboard + Alert Logic	OPEN / PENDING / RESOLVED / DISMISSED Money at risk vs. money protected	Audit trail; CFO sign-off as legal compliance document
5	Upcoding Detection	Billed amount > procedure maximum allowed procedure_rates lookup + overcharge%	Illegal — criminal fraud; carries fines. USD 68bn/yr (IJSRA 2024)
6	Phantom Billing / Ghost Patient	Patient for non-existent patient or no appointment on record for date	Pure fraud — Michigan case: \$61M Medicare fraud via ghost patients
7	Unbundling Detection	Bundled procedure codes split into multiple separate codes to inflate total	Illegal reimbursement inflation; appendectomy case: +€265 per episode
8	Credit Balance Accumulation	Total paid > total billed Days held vs. 180-day legal threshold	Government reporting required at 180 days in most EU jurisdictions
9	Surprise Billing / Out-of-Network	Block outside patient's insurance network bills at full private rates	Most contested issue in EU/US healthcare policy; patient never consents
10	Unregistered Doctor Detection	Doctor not registered in ANY insurance network (beyond simple network mismatch)	CRITICAL: potential fraud or unlicensed practitioner — verify credentials
11	Export to Dated Report File	Daily .txt compliance report with CFO sign-off block Atomic transaction-safe write	Legal compliance document; creates permanent daily audit archive
12	Excel Input/Output (pandas)	Read hospital_invoices.xlsx → scan → write HealthLedger_Alerts.xlsx	Option 1 integration: zero IT project; hospital sends/receives Excel file

B — Cooperation Agreement AI Compliance Layer

A RAG-based compliance engine trained on hospital cooperation agreements, delivering clause-traceable accounting determinations via a persistent managed AI agent (claude-opus-4-8). Addresses the 200+ cooperation agreement compliance problem — real-time guidance at point of accounting decision.

C — Fault Tolerance Architecture (Phase 2 — designed 28 March 2026)

- **Read-Only Access Enforcement:** All hospital connections opened in read-only mode. SELECT-only DB credentials. Hard-coded prohibition on any write to source data.
- **Progressive Scan with Checkpoint/Resume:** Scan results saved every 100 invoices (configurable). On crash, next run resumes from last checkpoint — never loses partial work.
- **Transaction-Safe File Writes:** Write to .tmp → verify → atomic rename. Prevents partial output files.
- **SHA-256 Data Integrity Hashing:** Source data hash computed before and verified after each scan. Unexpected modification → system halt + CRITICAL alert.
- **Automated Logging:** Every action, every invoice, every error timestamped in daily log file. Hospital sends log file for remote diagnosis — no IT knowledge required.
- **Error Escalation Protocol:** MINOR → silent log; MEDIUM → email notification; CRITICAL → immediate alert with full log attached.
- **Time-Limited Remote Access:** One-time session key per support call. Auto-expiry. All remote actions logged. No permanent backdoor access.
- **Installation Status Dashboard:** RUNNING_NORMAL / WARNING / OFFLINE across all hospital installations.

D — On-Site Immersion Methodology

The Custom Rules Engine — the configurable layer for institution-specific cooperation agreements — is intentionally built through 2–3 day on-site discovery sessions per hospital client. The inventor observes workflows, documents agreements, and codes custom rules from direct immersion. This creates the competitive moat: each hospital's rules are unique; switching to a competitor requires re-entering everything. The on-site methodology is an independently patentable inventive contribution.

5. DETAILED TECHNICAL DESCRIPTION

5.1 System Architecture — Five-Layer Model

Layer 1 — Hospital Data Interface

Read-only connection to hospital billing system or Excel import. SHA-256 hash computed on load. Invoices parsed into typed InvoiceRecord objects.

Layer 2 — Detection Engine

12 modular detection functions run sequentially over the invoice dataset. Each returns zero or more AlertRecord objects. Scan runs with checkpoint/resume fault tolerance (BATCH_SIZE=100).

Layer 3 — AI Agent Kernel

Persistent managed agent (Anthropic beta.agents.create, claude-opus-4-8) with institution-specific instructions. Receives augmented prompts from RAG engine; returns streamed compliance determinations.

Layer 4 — RAG Compliance Layer

Vector database of semantically indexed cooperation agreement clauses (6-type taxonomy). Semantic search on accounting query → top-5 clause retrieval → augmented prompt assembly.

Layer 5 — Output & Frontend

Transaction-safe file write (HealthLedger_Alerts.xlsx + dated .txt report). CFO Dashboard aggregation. Next.js 14 frontend deployed to healthledgerai.com via static export.

5.2 12-Feature Detection Engine — Key Algorithms

Each of the 12 features implements a specific algorithmic approach optimised for hospital-scale invoice datasets (10,000–100,000+ invoices/month):

- **Duplicate Detection:** Set-based fingerprint (vendor|amount|date). $O(1)$ lookup. Handles 100,000 invoices without performance degradation.
- **Risk Scoring:** Threshold function — CRITICAL $\geq \text{€}50\text{k}$, HIGH $\geq \text{€}10\text{k}$, MEDIUM $\geq \text{€}1\text{k}$, LOW $< \text{€}1\text{k}$.
- **Identity Collision:** Hash map indexed by (firstName, lastName, dob). $O(n)$ build + $O(1)$ query. Blocks invoice and demands parental name verification (Greek secondary identifier).
- **Upcoding:** procedure_rates dictionary lookup. Overcharge percentage = $(\text{billed} - \text{maxAllowed}) / \text{maxAllowed} \times 100$.
- **Phantom Billing:** Double cross-reference: (1) patientId $\in$ registry; (2) invoiceDate $\in$ appointments[patientId]. Three outcomes: IN_NETWORK, NO_APPOINTMENT, GHOST_PATIENT.
- **Unbundling:** Group invoices by (patientId, date). For each group, test procedure code set against bundle_rules. Match ≥ 2 codes → unbundling alert + overcharge calculation.
- **Credit Balance:** $\text{sum}(\text{payments}) - \text{billed}$. If > 0 : days_held = today - invoice_date. Urgency: ≥ 180 days → CRITICAL, ≥ 150 → HIGH, else MEDIUM.

- **Surprise Billing:** Get all networks containing doctorId. If empty → UNREGISTERED (CRITICAL). If patientNetwork not in doctorNetworks → OUT_OF_NETWORK (HIGH).
- **CFO Dashboard:** Aggregate alerts by risk level. Sum moneyAtRisk (CRITICAL+HIGH) and moneyProtected (MEDIUM+LOW). Assign all alerts OPEN status for lifecycle tracking.
- **Excel I/O:** pandas df.read_excel() → to_dict('records') → all features → pd.DataFrame(results) → to_excel(). Output: Invoice ID, Vendor, Amount, Date, Alert Type, Risk, Overcharge%.

5.3 AI Agent & Session Architecture

Persistent agent entity created once per institution (agents.create, claude-opus-4-8). Agent carries institution-specific instructions: chart of accounts, naming conventions, interpretation precedents from on-site capture. Per-query: new session created (sessions.create), augmented prompt submitted, response streamed token-by-token (target first-token latency: <800ms). Agent state persists across all sessions; session state is ephemeral and isolated per query.

5.4 Cooperation Agreement RAG Pipeline

- Step 1: Parse cooperation agreement PDFs → plain text (section numbers preserved as metadata)
- Step 2: Segment into individual clauses (NLP boundary detection)
- Step 3: Classify each clause into one of 6 types: RevenueRecognition, CostAllocation, ComplianceTrigger, RevenueThreshold, PenaltyClause, ReconciliationRule
- Step 4: Generate 1536-dimensional vector embedding per clause
- Step 5: Index embeddings in vector DB (Pinecone / pgvector) with metadata (clauseId, agreementId, type, date)
- Step 6: On accounting query: embed query → ANN search → top-5 clauses → augmented prompt → agent session → streamed clause-cited response → audit log

5.5 Domain Training & On-Site Immersion (Custom Rules Engine)

The Custom Rules Engine is built per institution through 2–3 day on-site discovery sessions. The inventor embeds with the hospital accounting team, observes real decisions, documents cooperation agreements, and codes institution-specific rules. This produces: (1) a chart of accounts mapping; (2) a naming convention dictionary; (3) a set of InterpretationRecord objects encoding team-specific precedents; (4) procedure_rates calibrated to national healthcare fee schedules (Cyprus and Greece). All of this is assembled into the agent instruction set via buildInstructions(). Updates after deployment require no model retraining — instruction refresh only.

5.6 Fault Tolerance & Remote Access Architecture (Phase 2 — 28 March 2026)

Designed for Phase 2 — live API connection to hospital information systems. Seven requirements must be met before any live hospital deployment:

- **R1 — Read-Only Access:** Hospital connections use SELECT-only DB credentials. Any connection not in READ_ONLY mode throws a security exception immediately. HealthLedger AI never writes to, updates, or deletes source hospital data under any circumstances.
- **R2 — Progressive Scan with Checkpoint/Resume:** Scan saves a checkpoint JSON file after every batch of 100 invoices. File written via transaction-safe atomic rename. On crash: next run

reads checkpoint, resumes from lastProcessedId. Accumulated alerts preserved. Checkpoint deleted on successful scan completion.

■ **R3 — Transaction-Safe Writes:** All output file writes: write to .tmp → verify content === written → fs.renameSync (atomic). If verification fails, .tmp is deleted and error thrown. Target and temp files are never in a partially-written state.

■ **R4 — SHA-256 Data Integrity Hashing:** Before scan begins: hash = computeDataHash(invoices). After scan: recompute and compare. If hashes differ (source modified during scan): CRITICAL alert sent, system halts, full audit triggered.

■ **R5 — Automated Logging:** Winston logger writes timestamped entries for every action, every invoice ID processed, every error with stack trace. Daily log file named healthledger_YYYY-MM-DD.log. Hospital sends log file for remote diagnosis — no technical knowledge required on hospital side.

■ **R6 — Error Escalation Protocol:** Three severity levels: MINOR → silent log only. MEDIUM → email to dev@healthledgerai.com with error description. CRITICAL → immediate email alert with full log file attached. Covers: system halt, data access failure, integrity violation, GHOST_PATIENT detection, UNREGISTERED_DOCTOR.

■ **R7 — Time-Limited Remote Access:** For errors requiring remote diagnosis: hospital admin activates a one-time session key (crypto.randomBytes(32)). Key expires after configurable duration (default: 2 hours). Remote team connects with READ + EXECUTE_DEBUGGER permissions only (never WRITE). All remote actions logged to audit trail. Session auto-closes on expiry — no permanent access granted.

5.7 Frontend Interface & Deployment

Next.js 14 static export (output: 'export') deployed to healthledgerai.com via GitHub Pages CNAME. React 18 client components. framer-motion 11 spring-physics animations. Cormorant Garamond + Jost typography (Google Fonts). Key components: IntroOverlay (4-phase particle burst), HeroScene (3D parallax), ParticleField (constellation physics), CustomCursor (dual-spring), Botanical (SVG lily), ProblemHook (sequential problem display), SectionHow (3-step method), SectionContact (research programme enrolment to info@healthledgerai.com).

5.8 Visual Design System

Defined in app/globals.css as CSS custom properties. Distinctive design system: warm parchment background (#F4EFE4), antique gold accent (#A08558), deep sage (#4A6B5A), SVG fractalNoise grain texture at 4% opacity, three-radial-gradient ambient layer, CSS 3D perspective (1200px), 3D word-flip headline animation (rotateX 90°→0°).

6. SOURCE CODE LISTINGS

All listings © 2026 HealthLedger AI — Maria Polychroniadou. Production code reproduced verbatim from the repository (mrply97/AI, GitHub). The detection engine and fault-tolerance modules are the upgraded production-grade implementations; not the original learning-stage code.

6.1 Core AI Agent Module

Listing 6.1 — managed-agents-starter.ts (production, verbatim)

[illegible]

```

if (!agentId) { await createAgent(); return; }

// Example augmented prompt (assembled by RAG engine – see Listing 5.3):
const prompt = `
RETRIEVED CLAUSES – Agreement #47 (MedTech GmbH, eff. 2025-01-01):
[AGR-47 §3.2]: Revenue from device maintenance shall be recognised in the
period in which the service obligation is fulfilled.
[AGR-47 §3.5]: Quarterly reconciliation required within 30 days of close.

SCAN RESULTS – Invoice batch 2026-Q3:
INV-4471: MedTech device maintenance, €13,241 – exceeds max rate €900 (1371% over)
INV-4472: Patient P-002 (Sofia Kosta) – IDENTITY COLLISION with P-008
INV-4473: Dr. Stavrakis – UNREGISTERED in all insurance networks

Provide compliance determination citing exact clauses. Use account code 4300.
`;

await runSession(agentId, prompt);
}
main().catch(console.error);

```

Listing 6.2 — lib/detection-engine.ts (production, complete)

```
import { InvoiceRecord, PatientRecord, DoctorRecord, AlertRecord } from "../types";
import { date, differenceInDays } from "date-fns";
```

[illegible]

[illegible]

```

invoices: InvoiceRecord[],
bundleRules: Record<string, { codes: Set<string>; bundledMax: number }>
): AlertRecord[] {
// Group by patient + date
const groups = new Map<string, InvoiceRecord[]>();
for (const inv of invoices) {
const key = `${inv.patientId}|${inv.date}`;
const group = groups.get(key) ?? [];
group.push(inv);
groups.set(key, group);
}
const alerts: AlertRecord[] = [];
for (const [, group] of groups) {
const procedureCodes = new Set(group.map(i => i.procedureCode));
for (const [bundleName, rule] of Object.entries(bundleRules)) {
const matched = [...rule.codes].filter(c => procedureCodes.has(c));
if (matched.length >= 2) {
const totalBilled = group.reduce((s, i) => s + i.amount, 0);
const overcharge = totalBilled - rule.bundledMax;
alerts.push({
invoiceId: group.map(i => i.id).join(", "),
type: "UNBUNDLING", risk: getRisk(overcharge),
message: `${bundleName}: ${matched.length} codes billed separately ` +
` (€${totalBilled} vs bundled max €${rule.bundledMax}, +€${overcharge})`,
});
}
}
}
return alerts;
}

```

[illegible]

6.3 Fault Tolerance & Remote Access Architecture

Listing 6.3 — lib/fault-tolerance.ts (Phase 2 architecture — 28 March 2026)

[illegible]

[illegible]

```
// Log file sent by hospital when support is needed – no technical knowledge required.

const logger = createLogger({
format: format.combine(
format.timestamp({ format: "YYYY-MM-DD HH:mm:ss" })),
format.printf(({ timestamp, level, message, ...meta }) =>
`${timestamp} [${level.toUpperCase()}] ${message} ${JSON.stringify(meta)}`)
),
transports: [
new transports.File({ filename: `./logs/healthledger_${new Date().toISOString().split("T")[0]}.log`
}),
new transports.Console({ level: "error" }),
],
});

// ■■■ REQUIREMENT 6: ERROR NOTIFICATION + ESCALATION PROTOCOL ■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■
// Severity levels: MINOR (silent log) → MEDIUM (email notification) → CRITICAL (immediate alert)

enum ErrorSeverity { MINOR, MEDIUM, CRITICAL }

interface SystemAlert {
hospitalId: string;
severity: ErrorSeverity;
errorCode: string;
message: string;
logFilePath: string;
timestamp: string;
}

async function escalateError(alert: SystemAlert): Promise<void> {
logger.error(`[${alert.errorCode}] ${alert.message}`, { hospitalId: alert.hospitalId });

if (alert.severity === ErrorSeverity.MINOR) {
return; // Silent log only
}

if (alert.severity === ErrorSeverity.MEDIUM) {
await sendEmailNotification({
to: "dev@healthledgerai.com",
subject: `HealthLedger AI – Warning at ${alert.hospitalId}`,
body: `Error: ${alert.message}\nLog: ${alert.logFilePath}`,
});
return;
}

if (alert.severity === ErrorSeverity.CRITICAL) {
// Immediate alert with full log attached
await sendCriticalAlert(alert.errorCode, alert.message);
await sendEmailNotification({
to: "dev@healthledgerai.com",
subject: `CRITICAL – HealthLedger AI HALT at ${alert.hospitalId}`,
body: `SYSTEM HALTED\nError: ${alert.message}\nattachment: alert.logFilePath`,
});
}
}

// ■■■ REQUIREMENT 7: TIME-LIMITED REMOTE ACCESS ■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■■
// Hospital admin activates a time-limited session key for remote support.
// Session closes automatically – no permanent backdoor access permitted.

interface RemoteAccessSession {
```

[illegible]

6.4 RAG Pipeline

Listing 6.4 — lib/ingestion-pipeline.ts (production)

```
// FILE: lib/ingestion-pipeline.ts | © 2026 HealthLedger AI – Maria Polychroniadou
// DESCRIPTION: Cooperation Agreement Ingestion & RAG Compliance Pipeline

interface AgreementClause {
  clauseId: string; // e.g. "AGR-47-§3.2"
  agreementId: string;
  text: string;
  type: ClauseType;
  obligations: string[];
  effectiveDate: Date;
  embedding?: number[]; // 1536-dim vector
}

enum ClauseType {
  RevenueRecognition = "revenue",
  CostAllocation = "cost",
  ComplianceTrigger = "compliance",
  RevenueThreshold = "threshold",
  PenaltyClause = "penalty",
  ReconciliationRule = "reconciliation",
}

// STEP 1 – Ingest agreement (called once per document)
async function ingestAgreement(pdfPath: string, agreementId: string) {
  const rawText = await parsePDF(pdfPath);
  const clauses = await segmentClauses(rawText);
  const tagged = await tagClauseTypes(clauses);
  const embedded = await embedClauses(tagged);
  await vectorDB.upsert(embedded, { agreementId });
  return embedded;
}

// STEP 2 – RAG compliance query (called per accounting question)
async function queryCompliance(agentId: string, userQuery: string) {
  const qEmbed = await embedQuery(userQuery);
  const topK = await vectorDB.search(qEmbed, { topK: 5 });
  const context = topK.map(c =>
    `${c.agreementId} ${c.clauseId}: ${c.text}`
  ).join("\n\n");
  const prompt = `Context:\n${context}\n\nQuestion: ${userQuery}\nCite clause.`;
  const response = await runSession(agentId, prompt);
  return {
    determination: response.text,
    sourceClauses: topK.map(c => c.clauseId),
    agreements: topK.map(c => c.agreementId),
    timestamp: new Date(),
  };
}
```

6.5 Domain Training Module

Listing 6.5 — lib/domain-training.ts (production)

```
// FILE: lib/domain-training.ts | © 2026 HealthLedger AI – Maria Polychroniadou
// DESCRIPTION: Institutional Knowledge Module + On-Site Immersion Data Capture

interface InstitutionProfile {
  hospitalId: string;
  name: string;
```

```

chartOfAccounts: { code: string; name: string }[];
namingConventions: Record<string, string>;
interpretations: InterpretationRecord[];
}

interface InterpretationRecord {
  clauseRef: string; // e.g. "AGR-47-§3.2"
  decision: string; // "Recognise in Q3, account 4300"
  rationale: string; // Expert justification from on-site session
  capturedBy: string;
  capturedAt: Date;
}

// Capture new on-site decision → immediately refreshes agent instructions
async function captureInterpretation(hospitalId: string, record: InterpretationRecord) {
  await db.interpretations.insert({ hospitalId, ...record });
  await refreshAgentInstructions(hospitalId); // No retraining required
}

async function buildInstructions(hospitalId: string): Promise<string> {
  const p = await db.institutions.findById(hospitalId);
  const h = await db.interpretations.findByHospital(hospitalId, { limit: 30 });
  return `
You are HealthLedger AI for ${p.name}.

CHART OF ACCOUNTS:
${p.chartOfAccounts.map(a => ` ${a.code}: ${a.name}`).join("\n")}

NAMING CONVENTIONS:
${Object.entries(p.namingConventions).map(([k,v]) => ` "${k}" = "${v}"`).join("\n")}

PRECEDENT DECISIONS (on-site captured):
${h.map(r => ` [${r.clauseRef}] → ${r.decision} (${r.rationale})`).join("\n")}

Rules: cite exact clause ID. Use institution account codes. Flag ambiguous clauses.
`;
}

```

6.6 Frontend + Deployment Config

Listing 6.6 — app/page.tsx + next.config.js + CNAME (production)

```

// FILE: app/page.tsx | © 2026 HealthLedger AI — Maria Polychroniadou
// Framework: Next.js 14 · React 18 · TypeScript · framer-motion 11

"use client";
// 12-Feature system pillars displayed in UI
const PILLARS = [
  "Cooperation Agreement Mapping",
  "Accounting Compliance",
  "Hospital-Specific AI",
  "On-Site Methodology",
];

// Problem statement — distilled from PhD research + European Interbalkan Medical Centre incident
const PROBLEM_HOOKS = [
  "A hospital has 200+ cooperation agreements.",
  "Each interpreted differently by accounting.",
  "One clause. Thousands in exposure.",
];

// Three-step method (maps to ingestion → domain training → live compliance)

```

```
const STEPS = [
  { num: "01", title: "Agreement Ingestion",
    body: "Map every clause – identify obligations, revenue thresholds, compliance triggers." },
  { num: "02", title: "Domain Training",
    body: "AI learns the institution's chart of accounts, naming conventions, and team practices." },
  { num: "03", title: "Live Compliance",
    body: "Real-time guidance. Answer is immediate – traceable to the source clause." },
];

// Brand identity constants
const BRAND_NAME = "HealthLedger AI";
const BRAND_DOMAIN = "healthledgerai.com"; // CNAME confirmed in repository
const BRAND_EMAIL = "info@healthledgerai.com";
const NAV_SUBTITLE = "Healthcare Accounting Intelligence";
const FOOTER_TAG = "Something important is being built.";
const COPYRIGHT = "© 2026 HealthLedger AI · All rights reserved";

// FILE: next.config.js – static export deployed to healthledgerai.com
// output: 'export' | trailingSlash: true | images: { unoptimized: true }

// FILE: CNAME – domain record confirming healthledgerai.com ownership
// healthledgerai.com
```

7. PRIOR ART ANALYSIS

General-Purpose AI (ChatGPT, Gemini, Claude default)

Distinction: Not trained on hospital cooperation agreements or billing fraud patterns. Cannot produce clause-cited auditable determinations. No fault-tolerant hospital data access architecture.

Legal Contract AI (Harvey AI, Luminance, Kira Systems)

Distinction: Focus on legal contract review, not accounting compliance. No billing fraud detection. No hospital chart-of-accounts integration. No on-site immersion methodology.

Healthcare ERP (SAP S/4HANA Healthcare, Oracle Health)

Distinction: Structured workflow systems requiring IT configuration of all rules. No natural-language compliance reasoning. No cooperation agreement text ingestion. No AI fraud detection.

Billing Audit Software (nThrive, Optum360)

Distinction: Rule-based detection without AI/ML reasoning. No cooperation agreement RAG layer. No persistent institutional knowledge agent. No on-site discovery methodology.

General RAG Systems

Distinction: The present invention distinguishes through: (a) 6-type healthcare accounting clause taxonomy; (b) institution-specific managed agent architecture; (c) on-site immersion methodology; (d) 12-feature fraud detection pipeline; (e) fault-tolerant read-only hospital deployment.

8. DRAFT PATENT CLAIMS (16 CLAIMS)

16 draft claims are presented: 5 independent (system, method, fault architecture, methodology, UI) and 11 dependent. Conception date: 14 March 2026. Final scope and language to be determined with counsel.

Claim 1 — Independent: Billing Fraud Detection System

A computer-implemented system for detecting billing fraud and accounting compliance violations in hospital financial operations, the system comprising: (a) a data ingestion module configured to receive hospital invoice records from at least one of: a hospital information system accessed in read-only mode, or an Excel spreadsheet file; (b) a duplicate invoice detection module using a set-based fingerprint comprising vendor identifier, invoice amount, and invoice date, to identify duplicate submissions with O(1) lookup time; (c) a risk scoring module classifying each detected alert as CRITICAL, HIGH, MEDIUM, or LOW based on the financial magnitude of the associated invoice amount; (d) a patient identity collision detection module cross-referencing invoice patient identifiers against a patient registry to detect patients sharing identical first name, last name, and date of birth, and blocking invoice processing pending parental name verification; (e) an upcoding detection module comparing billed procedure amounts against a procedure rates dictionary to identify billings exceeding the maximum allowed rate, and computing an overcharge percentage; (f) a phantom billing detection module cross-referencing each invoice against a patient registry and appointment records to classify invoices as IN_NETWORK, NO_APPOINTMENT, or GHOST_PATIENT; (g) an unbundling detection module grouping invoices by patient identifier and date, and testing each group against a bundle rules dictionary to detect deliberate procedure code splitting; (h) a credit balance detection module identifying accounts where total payments exceed total billed amount, computing days held, and applying a 180-day legal compliance threshold; (i) a surprise billing detection module classifying each billing doctor as IN_NETWORK, OUT_OF_NETWORK, or UNREGISTERED across all insurance networks; (j) a CFO dashboard aggregation module producing a structured report comprising alert counts by risk level, total money at risk, total money protected, and an alert lifecycle audit trail with OPEN, PENDING, RESOLVED, and DISMISSED states; (k) an Excel output module writing detection results to a structured spreadsheet file with columns for Invoice ID, Vendor, Amount, Date, Alert Type, Risk Level, and Overcharge Percentage.

Claim 2 — Independent: Fault Tolerance Architecture

A computer-implemented fault-tolerant architecture for operating a billing fraud detection system on live hospital information systems, the architecture comprising: (a) a read-only access enforcement module configured to open all hospital system connections with SELECT-only database credentials and to throw a security exception upon any attempt to open a non-read-only connection; (b) a progressive scan module processing invoice records in configurable batches and writing a checkpoint record after each batch, said checkpoint comprising the last successfully processed invoice identifier and all alerts accrued to that point, wherein a subsequent scan run detects an existing checkpoint and resumes processing from the recorded position; (c) a transaction-safe write module implementing an atomic write pattern comprising: writing content to a temporary file; verifying the written content matches the intended content; and performing an atomic rename of the temporary file to the target path, wherein partial writes are prevented; (d) a data integrity module computing a cryptographic hash of the source invoice dataset before the scan commences, recomputing the hash after scan completion, and halting the system with a CRITICAL alert if the hashes do not match; (e) an automated logging module recording every scan action, every processed invoice identifier, and every error with timestamp, severity, and full context to a dated log file; (f) an error escalation module applying a three-tier severity protocol wherein MINOR errors are silently logged, MEDIUM errors trigger an email notification to the development team, and CRITICAL errors trigger an immediate alert with the log file attached; and (g) a time-limited remote access module generating a cryptographically random one-time session key upon hospital administrator request, enforcing an automatic session expiry, restricting remote permissions to read and debugger-execute operations, and logging all remote actions to the audit trail.

Claim 3 — Independent: AI Compliance Method

A computer-implemented method for generating real-time accounting compliance guidance for a hospital institution, comprising: (a) ingesting the institution's cooperation agreements by parsing, clause segmentation, six-type obligation classification, vector embedding, and vector database indexing; (b) capturing institutional knowledge through on-site immersion, comprising observation of accounting decisions, structured knowledge elicitation, and recording of InterpretationRecord objects; (c) creating a persistent AI agent entity with institution-specific instructions assembled from the captured chart of accounts, naming conventions, and interpretation records; (d) on receipt of an accounting compliance query: embedding the query; retrieving top-K agreement clauses by semantic similarity; assembling a retrieval-augmented prompt citing clause identifiers; submitting to an agent session; streaming the compliance determination; and writing a structured audit log entry; and (e) updating the agent instruction set incrementally as new interpretation records are captured, without retraining the underlying language model.

Claim 4 — Independent: On-Site Immersion Methodology

A method for developing a domain-specific artificial intelligence compliance system for a healthcare institution, comprising: (a) conducting on-site engagement of 2–3 days with the institution's accounting department to observe real billing decisions and document cooperation agreements; (b) recording each observed decision as a structured InterpretationRecord comprising a clause reference, the decision taken, the expert rationale, and the date; (c) eliciting unwritten naming conventions and interpretive practices through structured interviews and formalising them as a named entity mapping dictionary; (d) calibrating a procedure_rates dictionary to the national healthcare fee schedule of the target jurisdiction; (e) assembling an AI agent instruction set from: the institution's chart of accounts; the named entity mapping; and curated InterpretationRecords as few-shot examples; (f) deploying the Custom Rules Engine with institution-specific rules configured during said on-site session; and (g) updating the instruction set and custom rules incrementally as new decisions are captured post-deployment, without model retraining.

Claim 5 — Independent: Animated Entry Experience Method

A computer-implemented method for an animated interface entry sequence comprising a four-phase diamond-glitter particle animation wherein: Phase 1 (0–2,600ms) renders a spherical particle swarm with Y-axis rotation and sinusoidal X-axis tilt; Phase 2 (2,600–5,400ms) expands the sphere radius by a factor of at least eight using easeInOutCubic interpolation; Phase 3 (5,400–6,400ms) holds the scattered particle field; and Phase 4 (6,400–7,800ms) transitions overlay opacity 1→0 using easeInQuad, revealing the underlying application interface.

Claim 6 — Dependent on Claim 1

The system of Claim 1, wherein the patient identity collision detection module of step (d) implements a hash map indexed by (firstName, lastName, dob) tuples for $O(n)$ build time and $O(1)$ per-query collision lookup, and wherein the module outputs a BLOCKED status requiring verification of parental names as a secondary patient identifier before invoice processing may continue.

Claim 7 — Dependent on Claim 1

The system of Claim 1, wherein the unbundling detection module of step (g) groups invoices by the composite key (patientId, invoiceDate), applies a set intersection between each group's procedure codes and each bundle rule's code set, and raises an alert when two or more codes from a single bundle rule appear in the same group, computing the overcharge as the sum of individual invoice amounts minus the bundled maximum rate.

Claim 8 — Dependent on Claim 1

The system of Claim 1, wherein the credit balance detection module of step (h) computes the credit amount as the sum of all payments minus the billed amount, and wherein the alert risk level escalates to CRITICAL when the days held value reaches or exceeds 180 days, corresponding to the government reporting threshold applicable in the target jurisdiction.

Claim 9 — Dependent on Claim 1

The system of Claim 1, wherein the surprise billing detection module of step (i) queries all insurance networks for the presence of the billing doctor identifier, returning UNREGISTERED with CRITICAL risk when the doctor is absent from all networks, and OUT_OF_NETWORK with HIGH risk when the doctor is present in other networks but not in the patient's specific insurer network.

Claim 10 — Dependent on Claim 2

The fault-tolerant architecture of Claim 2, wherein the cryptographic hash of step (d) is computed using the SHA-256 algorithm over a serialised representation of the invoice dataset comprising invoice identifiers, vendor names, amounts, dates, and patient identifiers, and wherein any detected mismatch triggers immediate system halt and dispatch of a CRITICAL escalation alert with the integrity violation code.

Claim 11 — Dependent on Claim 2

The architecture of Claim 2, wherein the time-limited remote access module of step (g) generates the session key using a cryptographically secure random number generator producing at least 256 bits of entropy, and wherein the session automatically expires after a configurable duration not exceeding two hours, with the expiry enforced server-side regardless of client state.

Claim 12 — Dependent on Claim 3

The method of Claim 3, wherein the agreement clause classification of step (a) uses a six-type taxonomy comprising RevenueRecognition, CostAllocation, ComplianceTrigger, RevenueThreshold, PenaltyClause, and ReconciliationRule, and wherein each classified clause is stored with metadata comprising the clause identifier, parent agreement identifier, obligation type, and effective date.

Claim 13 — Dependent on Claim 3

The method of Claim 3, wherein the AI agent entity is created using the Anthropic Managed Agents API by calling `beta.agents.create()` with the model identifier `claude-opus-4-8` and the institution-specific instruction set, and wherein the resulting agent identifier is stored durably for reuse across all subsequent compliance sessions without recreation.

Claim 14 — Dependent on Claim 4

The methodology of Claim 4, wherein the `procedure_rates` calibration of step (d) maps each procedure code to a tuple of (procedure name, minimum allowed fee, maximum allowed fee) sourced from the official national healthcare fee schedule of the target jurisdiction, and wherein said calibration is performed during the on-site engagement as a prerequisite to production deployment.

Claim 15 — Dependent on Claims 1 and 2

The system of Claim 1 implemented with the fault-tolerant architecture of Claim 2, wherein the progressive scan module of Claim 2 step (b) calls the detection modules of Claim 1 on each batch, and wherein the Excel output module of Claim 1 step (k) uses the transaction-safe write module of Claim 2 step (c) for all file write operations.

Claim 16 — Independent: Use

Use of the system of any one of Claims 1, 6, 7, 8, 9, or 15 for the purpose of reducing financial loss and legal liability in hospital accounting operations by automatically detecting billing fraud, identity errors, cooperation agreement compliance violations, and credit balance accumulation in hospital invoice datasets, and delivering structured alert reports to hospital CFOs and accounting departments.

9. ACADEMIC REFERENCES

The following peer-reviewed and professional sources were consulted during the development of HealthLedger AI and support the problem statement and non-obviousness arguments in this disclosure.

V■rzar, A. A. (2024). Assessing the Transformative Impact of AI Adoption on Efficiency, Fraud Detection, and Skill Dynamics in Accounting Practices. *Journal of Risk and Financial Management*, 17(12), 577.
<https://doi.org/10.3390/jrfm17120577>

Nimkar, S., et al. (2024). Increasing Trends of AI With Robotic Process Automation in Health Care. *Cureus*, 16(9): e69680.
<https://doi.org/10.7759/cureus.69680>

International Journal of Science and Research Archive (IJSRA). (2024). Fraud detection in healthcare billing and claims.
<https://doi.org/10.30574/ijrsra.2024.13.2.2606>

Sanz Martín, L., Parra-Domínguez, J., Corchado, J. M., & Zafra-Gómez, E. (2025). Recent evolution and growth of AI and advanced technologies in accounting and finance. *Spanish Journal of Finance and Accounting*.
<https://doi.org/10.1080/02102412.2025.2582120>

Becker's Hospital Review. (2024). 5 Common Accounting Blunders Hospitals Can Avoid.

Halunen Law. (2024). Upcoding and Unbundling Are Common Types of Healthcare Fraud.

Medical Billing Advocates of America. (2024). Up to 80% of medical bills contain at least one billing error. Cited in Locumtenens.com (2024).

Halunen Law / US DOJ. (2024). Michigan group fraudulently billed Medicare for \$61M using ghost patients (2023 case).

10. ABSTRACT

HealthLedger AI is an artificial intelligence system for billing fraud detection and accounting compliance in hospital finance departments. Conceived on 14 March 2026 and first reduced to practice on 17 March 2026, the system comprises: (i) a 12-feature billing fraud and error detection engine (duplicate invoices, risk scoring, patient identity collision, upcoding, phantom billing, unbundling, credit balance accumulation, surprise billing, unregistered doctor detection, CFO dashboard, and Excel input/output); (ii) a cooperation agreement compliance layer using retrieval-augmented generation over semantically indexed agreement clauses, delivered via a persistent managed AI agent (claude-opus-4-8) with institution-specific training; (iii) a fault-tolerant deployment architecture for live hospital systems comprising read-only access enforcement, progressive checkpoint/resume scanning, transaction-safe file writes, SHA-256 data integrity hashing, automated logging, three-tier error escalation, and time-limited remote access with cryptographic session keys; (iv) a novel on-site immersion methodology for capturing institutional accounting knowledge as AI training data, constituting the competitive moat of the Custom Rules Engine; and (v) a branded web interface deployed at healthledgerai.com with distinctive visual design, animated entry experience, and 3D parallax UI components. The brand name 'HealthLedger AI', domain healthledgerai.com, and associated taglines are claimed as trademarks. All source code and artwork are copyright © 2026 Maria Polychroniadou.

11. INVENTOR DECLARATION

I, the undersigned, declare that:

- I am the sole original inventor of all subject matter described in this disclosure. I conceived the HealthLedger AI system on Saturday 14 March 2026, beginning with zero prior coding knowledge. I achieved the first working reduction to practice on Tuesday 17 March 2026 (Version 1.0 — 11 detection features, Excel input/output). I designed the Phase 2 fault-tolerance architecture on 28 March 2026.
- The brand name 'HealthLedger AI', the domain healthledgerai.com, the email info@healthledgerai.com, and all taglines identified in Section A are my original creations and I claim all trademark rights therein.
- All source code, artwork, and design assets are original works of authorship created by me. Copyright subsists in my name as Maria Polychroniadou.
- All information in this disclosure is true and accurate to the best of my knowledge.
- I have not made any public disclosure of this invention that would prejudice patent filing. I will inform counsel immediately if any such disclosure occurs.
- I acknowledge my ongoing duty to disclose all prior art and related publications of which I am or become aware.

Full Legal Name: Maria Polychroniadou
Email: polychroniadou@gmx.de
Date of Conception: 14 March 2026
Date of Disclosure: 26 June 2026
Signature:

Witnessed By:

Witness Date:

Notary / Stamp:
