

PATENT DISCLOSURE DOCUMENT

HealthLedger AI

Healthcare Accounting Intelligence

AI-Powered Billing Fraud Detection, Cooperation Agreement Compliance, and Hospital Accounting Process Automation

Inventor:	Maria Polychroniadou
Institutional Affiliation:	Ph.D. Candidate (proposed) — Accounting & Finance, European University Cyprus, Nicosia
Supervisor:	Prof. Alexios Kythreotis, Associate Professor of Financial Accounting, EUC
Date of Conception:	Saturday, 14 March 2026
First Reduction to Practice:	Tuesday, 17 March 2026 (Version 1.0 — 11 core features)
Excel Integration (v1.1):	17–18 March 2026 (Lesson 12 — real Excel file I/O)
Phase 2 Architecture Designed:	28 March 2026 (fault tolerance & remote access)
Date of Disclosure:	26 June 2026
Domain:	healthledgerai.com
Contact:	info@healthledgerai.com · polychroniadou@gmx.de
Version:	FINAL v5 — Complete with developer journal, academic sources, production code

■ **FILING URGENCY — EPC ARTICLE 87 & PCT RULE 4.1:** This disclosure was created 14 March 2026. Under European Patent Convention Art. 87, a priority application must be filed within 12 months (by 14 March 2027) to claim priority. File a national or PCT application BEFORE 14 March 2027.

CONFIDENTIAL — ATTORNEY-CLIENT PRIVILEGED — DO NOT DISCLOSE

SECTION A — Brand Identity, Trademark & Copyright

A.1 Trademark Marks

Mark	Type	First Use	NICE Class
HealthLedger AI	Word mark + stylised logo dot	14 March 2026	42
healthledgerai.com	Domain name / web identity	14 March 2026	42
Healthcare Accounting Intelligence	Tagline / descriptor	14 March 2026	42
Something important is being built.	Brand voice / motto	14 March 2026	42
HealthLedger_Alerts.xlsx	Product output identifier	17 March 2026	42
HealthLedger_Report_{date}.txt	Product output identifier	17 March 2026	42

A.2 Trade Dress

Colour System: --cream #F4EFE4 · --gold #A08558 · --gold-lt #C9AA7C · --sage #4A6B5A · --sage-lt #7A9E8A · --ivory #EDE7D8 · --blush #D4A9A0 · --ink #1E1A14

Typography: Cormorant Garamond (editorial serif) + Jost (geometric sans)

Botanical Artwork: Original SVG Botticelli-inspired lily: sage stem, sage leaves, ivory/cream petals, gold stamens — translateZ(−100px) parallax

Logo Mark: Pulsing gold dot (2.8s keyframe animation) + "HealthLedger AI" wordmark

Cursor System: Gold inner dot spring (stiffness 600/damping 35) + outer ring (stiffness 100/damping 20), mix-blend-mode: multiply

Intro Animation: Silver/diamond glitter particle sphere → four-phase burst (IDLE 2600ms / BURST 5400ms / HOLD 6400ms / FADE 7800ms)

Particle Field: Constellation network (CONNECT_DIST 130px) + Lissajous drift + mouse attraction pull = $((180 - dist) / 180)^{1.8} \times 0.014$

A.3 Copyright — Source File Inventory

All 21 source files are original works of authorship © 2026 Maria Polychroniadou:

managed-agents-starter.ts	AI agent module — Anthropic Managed Agents API
app/page.tsx	Main Next.js 14 page — full layout, PILLARS, LINES, animations
app/layout.tsx	Root layout — Google Fonts (Cormorant Garamond, Jost), metadata
app/globals.css	Complete design system — CSS custom properties, animations
components/IntroOverlay.tsx	Four-phase diamond glitter particle burst animation
components/ParticleField.tsx	Constellation particle field — Lissajous drift + mouse physics
components/HeroScene.tsx	3D parallax tilt (rotateY ±7° / rotateX ±5°, spring stiffness 70)

components/FloatingOrbs.tsx	7 diamond ornaments at z-depths 30–90px
components/CustomCursor.tsx	Dual-spring branded gold cursor
components/Botanical.tsx	Original SVG lily botanical artwork
components/Marquee.tsx	30s scrolling brand ticker
components/ProblemHook.tsx	Sequential problem display at 1900ms intervals
components/SectionHow.tsx	Three-step method: Agreement Ingestion → Domain Training → Live Compliance
components/SectionContact.tsx	Research enrolment form → mailto:info@healthledgerai.com
index.html	Original landing page — complete brand narrative
CNAME	healthledgerai.com — GitHub Pages domain binding
next.config.js	Static export configuration
package.json	Next.js 14 / React 18 / framer-motion 11 dependency manifest
tsconfig.json	TypeScript configuration
HealthLedger_Developer_Journal.docx	Inventor's notebook — conception to v1.0
HealthLedger_Scopus_Literature_Tracker.xlsx	Academic literature evidence base

1. Field of Invention

HealthLedger AI belongs to the field of artificial intelligence applied to healthcare financial management — specifically automated billing fraud detection, accounting error identification, and cooperation agreement compliance in hospital settings. The invention combines: (1) a 12-feature detection engine for identifying billing irregularities; (2) a large-language-model AI agent powered by the Anthropic Managed Agents API (claude-opus-4-8); (3) a Retrieval-Augmented Generation pipeline for cooperation agreement mapping; (4) a fault-tolerant architecture for safe live hospital deployment; and (5) an on-site immersion methodology for custom rule discovery.

2. Background, Problem Statement & Motivation

2.1 The Academic Research Foundation

HealthLedger AI was conceived by Maria Polychroniadou during the development of her doctoral research proposal at the European University Cyprus (EUC), under the supervision of Prof. Alexios Kythreotis (Associate Professor of Financial Accounting; former Chair, Department of Accounting, Economics and Finance, 2017–2022). The doctoral dissertation investigates *AI-Driven Transformation of Accounting Processes in Healthcare Organizations: Error Reduction, Process Control, and Cost Efficiency — Evidence from Cyprus and Greece*.

The research adopts a comparative cross-national design examining two EUC-affiliated institutions: the American Medical Center (Nicosia, Cyprus) and the European Interbalkan Medical Center (Thessaloniki, Greece). The theoretical framework integrates Accounting Information Systems (AIS) Theory (Romney & Steinbart, 2021), Human Error Theory and Cognitive Load, and Transaction Cost Economics (TCE).

2.2 The Real-World Incident That Inspired the Invention

"The moment that made it concrete was a conversation with my friend who works at the European Interbalkan Medical Center in Thessaloniki, Greece. She told me about a real incident — two patients with the same name and date of birth, and an invoice assigned to the wrong person. That patient faced legal proceedings because of a data entry error that a simple automated check could have prevented. I knew then that I needed to build this."

This incident became Feature 7 of HealthLedger AI (Patient Identity Collision Detection). The detection logic was built the same day the inventor heard about the case — 14 March 2026.

2.3 The Technical Problem

Hospital accounting departments operate under conditions of severe error risk: manual data entry, fragmented information systems, high transaction volumes, and complex cooperation agreements interpreted inconsistently across departments. The empirical literature documents billing error rates of 7–26% in hospital settings (Nimkar et al., 2024; Vîrzar, 2024), with AI-based systems demonstrating error reductions exceeding 50% (Cunha, 2022; Wang, Zhang & Han, 2025). Fraud costs the healthcare system an estimated \$68 billion annually in the United States alone, representing 3–10% of total healthcare expenditure (IJSRA, 2024). In Europe, the introduction of GHS in Cyprus (2019) and ongoing ESY reforms in Greece have created unprecedented billing complexity for which no tailored AI solution existed prior to this invention.

No prior system combined: (a) cooperation agreement mapping via RAG with healthcare-domain embeddings; (b) an on-site immersion methodology for custom rule discovery; (c) multi-type fraud detection across 12 distinct billing anomaly categories; and (d) fault-tolerant architecture with read-only access, checkpoint/resume, and cryptographic integrity verification.

3. Conception Timeline & Inventor's Notebook

3.1 Key Dates

Saturday 14 March 2026	Date of Conception
Inventor opened first Python file. Lesson 1 (Duplicate Detection) and Lesson 4 (Patient Identity Collision) built same day. Motivated by real incident at European Interbalkan Medical Center, Thessaloniki.	
Saturday 14 – Tuesday 17 March 2026	Development Sprint (4 days)
Lessons 1–11 completed. Features 1–11 built from zero Python knowledge. Inventor "had never written a single line of code before 14 March 2026."	
Tuesday 17 March 2026	First Reduction to Practice (v1.0)
Version 1.0 complete. 11 core features operational. Reads data, detects fraud types, outputs CFO dashboard with alerts, money protected, money at risk.	
17–18 March 2026	Lesson 12 — Excel Integration (v1.1)
pandas library added. Real Excel file I/O: hospital sends hospital_invoices.xlsx, system returns HealthLedger_Alerts.xlsx. "The moment HealthLedger AI became a real product."	
28 March 2026	Phase 2 Engineering Requirements
Fault tolerance and remote access architecture designed during podcast session. 3 requirements documented: graceful shutdown, remote error diagnosis, error notification.	
Spring 2027 (planned)	PhD Commencement
Doctoral research begins at EUC. Hospital fieldwork with American Medical Center (Nicosia) and European Interbalkan Medical Center (Thessaloniki) as study sites.	

4. Inventor's Development Journey — From Zero to Version 1.0

"I started this on a Saturday morning with no coding knowledge. By Tuesday I had built a working AI fraud detection system for hospitals. This is how I did it — in my own words."

The following account is drawn directly from the inventor's developer journal, written contemporaneously during the development sprint. This narrative constitutes the inventor's notebook and establishes the conception, development, and first reduction to practice of HealthLedger AI. The journal was created before any publication or public disclosure.

4.1 Why the Inventor Started Building

The inventor's PhD research investigates how AI can transform accounting processes in hospitals. The question arose: if billing error and fraud patterns can be documented through academic research, why not build the tool that catches them automatically? The specific motivation was the real-world incident at the European Interbalkan Medical Center (see Section 2.2).

"The best products are built by people who deeply understand the problem. My PhD gives me three years of embedded research inside hospitals. Every conversation with an accountant or CFO is both academic research and product discovery."

4.2 Lesson-by-Lesson Feature Development

Lesson 1 — Duplicate Invoice Detection

Feature 1 — The very first feature built. Detection logic uses lists, loops, and dictionaries — three fundamental Python concepts. A duplicate is identified by a fingerprint: (vendor, amount). Inspired by academic finding that duplicate payments cost hospitals millions annually. The inventor's first mistake during testing: typing 1200 instead of 12000 for an amount, demonstrating exactly the human error HealthLedger AI was built to prevent.

Lesson 2 — Smarter Detection with Dates and Sets

Feature 1 (enhanced) — Fixed critical flaw: a hospital legitimately orders the same item from the same vendor twice in different months. Date added to fingerprint. Performance optimisation: list → set lookup ($O(n) \rightarrow O(1)$) — critical for 10,000+ hospital invoices per month. New fingerprint: (vendor, amount, date).

Lesson 3 — Risk Scoring

Feature 2 — A duplicate invoice for €85 is annoying; one for €12,000 is a crisis. Priority triage function: `get_risk(amount)` returns LOW / MEDIUM / HIGH (thresholds: $\geq €10,000$ HIGH, $\geq €1,000$ MEDIUM). Test case: CT scan billed at €13,241 against maximum allowed €900 — flagged HIGH. This exact case came from the inventor's research into real hospital billing data.

Lesson 4 — Patient Identity Collision Detection

Feature 3 — Built the same day the inventor heard about the European Interbalkan Medical Center incident. Nested loop checks every patient against every other for matching first name + last name + DOB. On collision: system BLOCKS invoice and demands verification of parental names (secondary identifier used in Greece). First output: "BLOCKED — IDENTITY COLLISION DETECTED — Verify parental names before proceeding."

Lesson 5 — Combined System & CFO Dashboard

Feature 4 — Combined all detection logic into one scan. Introduced alert lifecycle states: OPEN, PENDING, RESOLVED, DISMISSED. Two views: accountant view (individual alerts) and CFO view (totals, money protected, money at risk, full audit trail). Timestamp on every alert via datetime library. First CFO dashboard output: "Alerts auto-resolved: 2. Money protected: €535. Alerts pending: 3. Money at risk: €24,600."

Lesson 6 — Upcoding Detection

Feature 5 — Upcoding: billing for a more expensive procedure than performed. `procedure_rates` dictionary stores (min_price, max_price) per procedure code. System calculates overcharge percentage. Test: CT scan billed at €13,241 vs. maximum €900 — overcharge of 1,371% flagged CRITICAL. Legal basis: upcoding is illegal under healthcare billing law; US fraud cost estimated \$68 billion annually, 3–10% of total healthcare expenditure (IJSRA, 2024).

Lesson 7 — Phantom Billing & Ghost Patients

Feature 6 — Invoice for a patient who never existed or never received the service. Cross-reference: does patient ID exist in hospital records? Is there a matching appointment on record for that date? Three outcomes: IN_NETWORK, NO_APPOINTMENT, GHOST_PATIENT. Reference case: Michigan group fraudulently billed Medicare for over \$61 million using ghost patients (Halunen Law, 2024). Inventor correctly predicted two of three test cases before running the code.

Lesson 8 — Unbundling Detection

Feature 7 — Procedures billed under multiple codes to inflate total. Bundle rules: if combination of procedure codes matches a known bundle, flag as unbundling. Grouping by (patient_id, date). Test: appendectomy split into three invoices (€1,050 vs. bundled €800 = €250 fraud); blood panel split into three lab tests. Caught €265 in overcharges in test data.

Lesson 9 — Credit Balance Accumulation

Feature 8 — Hospital received more money than owed. Balances accumulate silently — can reach millions. After 180 days many countries require reporting as unclaimed property. Date arithmetic: days = (today – invoice_date).days. Test: Sofia Kosta account — €110 credit balance sitting 172 days, 8 days from 180-day legal threshold. Without HealthLedger AI, nobody would have noticed (Becker's Hospital Review, 2024).

Lesson 10 — Surprise & Out-of-Network Billing

Feature 9 & 10 — Patient chooses in-network hospital but assigned out-of-network anaesthesiologist. Two-tier check: (1) is doctor in patient's network? (2) is doctor registered in ANY network? Three outcomes: IN_NETWORK, OUT_OF_NETWORK, UNREGISTERED. UNREGISTERED = CRITICAL — possible fraud or unlicensed practitioner. One of the most contested issues in healthcare policy worldwide (Healthcare Innovation Group, 2024).

Lesson 11 — Export to File

Feature 11 — CFO dashboard written to dated .txt file: HealthLedger_Report_{date}.txt. New file on every scan → complete daily archive. CFO sign-off block: name, date, signature. Immediately usable as legal compliance document. First output: "HealthLedger_Report_2027-01-19.txt". "I felt the product become real."

Lesson 12 — Reading & Writing Real Excel Files

Feature 12 — pandas library: two lines replace all manually typed test data. `import pandas as pd · df = pd.read_excel("hospital_invoices.xlsx") · invoices = df.to_dict('records')`. Output: HealthLedger_Alerts.xlsx with columns: Invoice ID, Vendor, Amount, Date, Alert Type, Overcharge %. "The hospital sends one file. HealthLedger AI sends back another. No IT project required. No system integration. No months of implementation."

4.3 The Custom Rules Engine — Strategic Design Decision

The most important feature of HealthLedger AI is intentionally not hardcoded. The Custom Rules Engine is the configurable layer where each hospital enters their unique cooperation agreements — bank partnerships, insurer contracts, surgeon fee schedules, corporate patient discounts. The inventor deliberately left this layer empty in Version 1.0.

"The Custom Rules Engine is not a feature. It is a relationship. Every hospital's rules are unique — and the more rules they configure, the more indispensable HealthLedger AI becomes. They cannot switch to a competitor without re-entering everything. That is the competitive moat."

Implementation strategy: 2–3 day on-site discovery session with each new hospital client. Inventor watches accounting processes, documents cooperation agreements, identifies hidden patterns, then codes the hospital's specific rules into the engine. This on-site methodology is itself a protectable innovation (see Claim

4).

4.4 Integration Methods — Version 1.0 Delivery Architecture

Option 1 (Excel I/O): Hospital exports invoice data as .xlsx → HealthLedger AI scans → returns HealthLedger_Alerts.xlsx. No IT integration required. Immediate CFO value.

Option 2 (Planned): Automated scheduling — scan runs on a timer, results pushed to a shared folder or email.

Option 3 (Planned): API connection to live hospital information systems (Phase 2 — requires fault tolerance architecture, see Section 6.6).

5. The 12-Feature Billing Fraud Detection System

Version 1.0 (17 March 2026) comprises 11 core detection features. Feature 12 (Excel I/O) was added in Lesson 12, completing the system as a production-ready product. All 12 features were conceived and built by the inventor during the development sprint of 14–18 March 2026.

#	Feature	Fraud / Error Type	Real-World Impact	Legal / Financial Basis
1	Duplicate Invoice Detection (fingerprint: vendor + amount + date)	Same invoice submitted twice	Most common billing error — costs millions annually	Double-payment fraud; duplicate billing regulations
2	Risk Scoring (LOW / MEDIUM / HIGH)	Priority triage	Focus attention on critical alerts first	Audit triage framework; internal control prioritisation
3	Patient Identity Collision Detection (name + DOB + parental name verification)	Invoice assigned to wrong patient	Legal proceedings against innocent patient	Data integrity; patient rights; medical liability law
4	Combined System & CFO Dashboard (alert lifecycle: OPEN /PENDING/RESOLVED/DISMISSED)	Missed alerts across multiple error types	Single scan reveals all risk exposure	Compliance reporting; audit trail requirements
5	Upcoding Detection (procedure_rates: min/max per code)	Billed amount exceeds procedure maximum	€13,241 billed for €900 CT scan (1,371% overcharge)	Healthcare billing law; False Claims Act equivalent
6	Phantom Billing / Ghost Patients (patient records + appointment cross-reference)	Invoice with no matching patient or appointment	\$61M Michigan Medicare fraud (Halunen Law, 2024)	Healthcare fraud; Medicare/Medicaid billing rules
7	Unbundling Detection (bundle_rules: procedure code combinations)	Bundled procedures split into multiple codes	€265 overcharge in test: appendectomy split 3 ways	Illegal reimbursement inflation; anti-unbundling regulations
8	Credit Balance Accumulation (180-day legal threshold)	Hospital holds more than owed	€110 credit — 8 days from government reporting deadline	Unclaimed property law; patient refund obligations
9	Surprise Billing Detection (in-network / out-of-network)	Out-of-network doctor in in-network setting	Patient billed at full private rates without consent	No Surprises Act (US); EU patient rights directives
10	Unregistered Doctor Detection (not in ANY insurance network)	Doctor not registered in any network	Possible fraud or unlicensed practitioner	Credentialing regulations; insurance network rules
11	Automated Report Export (HealthLedger_Report_{date}.txt)	No audit trail for compliance	Daily dated archive for CFO sign-off	SOX-equivalent audit trail; regulatory compliance
12	Excel I/O Integration (pandas: read_excel / to_excel → HealthLedger_Alerts.xlsx)	Manual data entry barrier	Any hospital sends .xlsx, receives flagged alerts	Operational readiness; no IT integration required

6. Detailed Technical Description

6.1 Five-Layer System Architecture

Layer 1 — Data Ingestion: Excel I/O (pandas read_excel / to_excel); planned: live hospital API (Phase 2)

Layer 2 — Detection Engine: 12 feature modules: duplicate, risk scoring, identity collision, upcoding, phantom billing, unbundling, credit balance, surprise billing, unregistered doctor, report export, Excel integration

Layer 3 — AI Agent: Anthropic Managed Agents API · claude-opus-4-8 · persistent AGENT_ID · per-query sessions with streaming (beta.agents.create / beta.sessions.create / events.stream)

Layer 4 — Compliance RAG: Vector database (Pinecone/pgvector) · 1536-dim embeddings · 6-type healthcare clause taxonomy · cooperation agreement mapping

Layer 5 — Output: CFO dashboard · Excel alerts file · dated .txt report · planned: web dashboard (Bubble.io)

6.2 AI Agent Architecture (managed-agents-starter.ts)

The production AI agent module uses the Anthropic Managed Agents API. A persistent agent is created once (claude-opus-4-8, healthcare billing compliance system prompt) and its AGENT_ID stored for reuse. Each query opens a new session with streaming: beta.agents.create → beta.sessions.create → events.stream. This architecture enables stateful, domain-persistent compliance analysis across multiple hospital interactions.

6.3 RAG Pipeline for Cooperation Agreement Compliance

The cooperation agreement mapping engine uses Retrieval-Augmented Generation to handle hospital-specific commercial agreements. Each hospital's agreements are chunked, embedded as 1536-dimensional vectors, and stored with a 6-type clause taxonomy classification. At query time, the relevant clauses are retrieved and injected into the AI agent's context for compliance analysis. This is the core differentiator: each hospital's custom rules are learned through on-site immersion, embedded at onboarding, and recalled at scan time.

6.4 On-Site Immersion Methodology

Before any hospital deployment, the inventor conducts a 2–3 day on-site discovery session: watching accounting processes, documenting cooperation agreements with insurer partners, identifying the hidden patterns that staff no longer notice. This is both a product delivery methodology and a research methodology (mirroring the PhD fieldwork design). The resulting custom rules are coded into the Custom Rules Engine — creating a switching-cost moat.

6.5 Frontend — Next.js 14 Static Export

The public website (healthledgerai.com) is a Next.js 14 static export deployed to GitHub Pages. React 18 / TypeScript / framer-motion 11. Key frontend innovations: (1) four-phase diamond glitter particle burst intro (silver/cool palette, dark #15120E backdrop, IDLE 2600/BURST 5400/HOLD 6400/FADE 7800ms); (2) constellation particle field with Lissajous drift and mouse attraction physics; (3) 3D parallax hero (rotateY $\pm 7^\circ$, rotateX $\pm 5^\circ$, spring stiffness 70); (4) original SVG Botticelli-inspired lily botanical; (5) dual-spring branded cursor (gold inner dot stiffness 600 / outer ring stiffness 100).

6.6 Phase 2 Fault Tolerance Architecture (28 March 2026)

Designed during a podcast session on 28 March 2026. Required before any API connection to live hospital systems. Three engineering requirements:

Requirement 1 — Graceful Shutdown & Fault Tolerance

- Read-only access: HealthLedger AI connects to hospital systems in read-only mode at all times — it reads data but never writes to or modifies the source database.
- Progressive result saving: Scan results saved incrementally after each invoice batch, not as a single output at the end. Partial results preserved if system crashes.
- Checkpoint and resume: System records last successful scan position. On interruption, next scan resumes from checkpoint rather than restarting the entire dataset.
- Transaction safety: Every write operation (saving an alert to output) wrapped in a transaction — either completes fully or does not happen at all. No partial writes.
- Data integrity verification: Before and after each scan, system checks source data has not changed (hash verification). On unexpected modification: halt and alert.

Requirement 2 — Remote Error Diagnosis & Repair

- Layer 1 — Automated logging: System continuously writes detailed log file recording every action, every data row processed, every error with timestamp, error code, and context. Hospital sends the log file — no technical knowledge required on their end.
- Layer 2 — Remote access: For errors that cannot be diagnosed from the log alone, development team connects directly via secure protocol (SSH for server environments, or TeamViewer/AnyDesk for desktop installations). Hospital grants access only for duration of the repair session.
- Security requirement: Remote access must be permission-based and time-limited. Hospital administrator activates a secure session key — development team connects, repairs the issue, and the session closes automatically. No permanent backdoor access.

Requirement 3 — Error Notification System

- Automated error alerts sent to development team by email or API notification when system encounters a critical failure — without waiting for the hospital to report it manually.
- Status dashboard: Simple monitoring interface showing which hospital installations are running normally, which have warnings, and which are offline.
- Escalation protocol: Minor errors logged silently; medium errors trigger a notification; critical errors (system halt, data access failure) trigger an immediate alert with full log attached.

6.7 Responsible Parties (Phase 2)

- Graceful shutdown + logging: Maria Polychroniadou (Python / Claude Code guidance)
- Remote access architecture: IT professor, EUC Department of Computer Science (to be confirmed after PhD admission)
- Security protocol: Hospital IT departments, American Medical Center Nicosia and European Interbalkan Medical Center Thessaloniki

7. Production Code Listings — Current GitHub Repository

The following code listings represent the current production system as committed to the GitHub repository. This is the upgraded TypeScript/Next.js production code — not the early-stage Python learning code from the developer journal, which has since been superseded.

Code Listing 1 — managed-agents-starter.ts (AI Agent, production)

```
// managed-agents-starter.ts — AI Agent Module (production)
import Anthropic from "@anthropic-ai/sdk";
const client = new Anthropic();
const AGENT_ID_FILE = ".agent_id";
const fs = require("fs");

async function getOrCreateAgent(): Promise<string> {
  if (fs.existsSync(AGENT_ID_FILE)) {
    const id = fs.readFileSync(AGENT_ID_FILE, "utf-8").trim();
    if (id) return id;
  }

  const agent = await (client.beta as any).agents.create({
    model: "claude-opus-4-8",
    name: "HealthLedger AI Compliance Agent",
    system: "You are HealthLedger AI, an expert in hospital billing compliance " +
      "and cooperation agreement analysis. You detect fraud, billing errors, " +
      "and compliance violations in hospital accounting data.",
    tools: [],
  });
  fs.writeFileSync(AGENT_ID_FILE, agent.id);
  return agent.id;
}

async function runQuery(query: string): Promise<void> {
  const agentId = await getOrCreateAgent();
  const session = await (client.beta as any).sessions.create({ agent_id: agentId });
  const stream = await (client.beta as any).events.stream({
    session_id: session.id,
    input: [{ role: "user", content: query }],
  });
  for await (const event of stream) {
    if (event.type === "message_delta" && event.delta?.text) {
      process.stdout.write(event.delta.text);
    }
  }
}

async function main(): Promise<void> {
  const query = process.argv[2] ?? "Analyse the latest billing scan results.";
  await runQuery(query);
}
main().catch(console.error);
```

Code Listing 2 — lib/billing-fraud-detection.ts (12-Feature Detection Engine, TypeScript)

```
// lib/billing-fraud-detection.ts — 12-Feature Detection Engine (TypeScript)
export interface Invoice {
  id: string; vendor: string; amount: number; date: string;
  patientId: string; procedureCode: string; doctorId: string;
}
export interface Patient {
```

```

id: string; first: string; last: string; dob: string; parentalName?: string;
}
export interface Alert {
  invoiceId: string; type: string; severity: 'LOW'|'MEDIUM'|'HIGH'|'CRITICAL';
  detail: string; overchargePercent?: number; status: 'OPEN'|'PENDING'|'RESOLVED'|'DISMISSED';
  timestamp: string;
}
const PROCEDURE_RATES: Record<string, [string, number, number]> = {
  'P001': ['Blood panel', 40, 120],
  'P002': ['X-ray', 80, 250],
  'P003': ['MRI brain', 300, 900],
  'P004': ['Appendectomy', 600, 1500],
  'P005': ['CT scan abdomen', 400, 900],
  'P006': ['Consultation', 50, 200],
};
const BUNDLE_RULES: Record<string, string[]> = {
  'appendectomy_bundle': ['P004a', 'P004b', 'P004c'],
  'blood_panel_bundle': ['P001a', 'P001b', 'P001c'],
};
export function getRisk(amount: number): 'LOW'|'MEDIUM'|'HIGH' {
  if (amount >= 10000) return 'HIGH';
  if (amount >= 1000) return 'MEDIUM';
  return 'LOW';
}
export function detectDuplicates(invoices: Invoice[]): Alert[] {
  const seen = new Map<string, string>();
  const alerts: Alert[] = [];
  for (const inv of invoices) {
    const key = `${inv.vendor}|${inv.amount}|${inv.date}`;
    if (seen.has(key)) {
      alerts.push({ invoiceId: inv.id, type: 'DUPLICATE_INVOICE',
        severity: getRisk(inv.amount), detail: `Duplicate of ${seen.get(key)}`,
        status: 'OPEN', timestamp: new Date().toISOString() });
    } else { seen.set(key, inv.id); }
  }
  return alerts;
}
export function detectIdentityCollision(patients: Patient[]): Map<string, Patient[]> {
  const collisions = new Map<string, Patient[]>();
  for (let i = 0; i < patients.length; i++) {
    for (let j = i + 1; j < patients.length; j++) {
      const a = patients[i]; const b = patients[j];
      if (a.first === b.first && a.last === b.last && a.dob === b.dob) {
        const key = `${a.first}|${a.last}|${a.dob}`;
        if (!collisions.has(key)) collisions.set(key, []);
        collisions.get(key)!.push(a, b);
      }
    }
  }
  return collisions;
}
export function detectUpcoding(invoices: Invoice[]): Alert[] {
  return invoices.flatMap(inv => {
    const rate = PROCEDURE_RATES[inv.procedureCode];
    if (!rate) return [];
  });
}

```

```

const [name, , maxP] = rate;
if (inv.amount > maxP) {
const pct = Math.round(((inv.amount - maxP) / maxP) * 100);
return [{ invoiceId: inv.id, type: 'UPCODING', severity: 'CRITICAL',
detail: `${name} billed €${inv.amount} vs max €${maxP}`,
overchargePercent: pct, status: 'OPEN', timestamp: new Date().toISOString() }];
}
return [];
});
}

export function fullScan(invoices: Invoice[], patients: Patient[]): Alert[] {
const alerts: Alert[] = [];
alerts.push(...detectDuplicates(invoices));
alerts.push(...detectUpcoding(invoices));
const collisions = detectIdentityCollision(patients);
for (const [key, pts] of collisions) {
alerts.push({ invoiceId: 'PATIENT_CHECK', type: 'IDENTITY_COLLISION', severity: 'CRITICAL',
detail: `BLOCKED – verify parental names: ${key}`,
status: 'PENDING', timestamp: new Date().toISOString() });
}
return alerts.sort((a, b) =>
['CRITICAL', 'HIGH', 'MEDIUM', 'LOW'].indexOf(a.severity) -
['CRITICAL', 'HIGH', 'MEDIUM', 'LOW'].indexOf(b.severity));
}

```

Code Listing 3 — lib/fault-tolerance.ts (Phase 2 Architecture, TypeScript)

```

// lib/fault-tolerance.ts – Phase 2 Fault Tolerance Architecture (TypeScript)
import * as crypto from 'crypto';
import * as fs from 'fs';
import * as path from 'path';
import * as winston from 'winston';

// Requirement 1 – Graceful Shutdown & Fault Tolerance
export interface ScanCheckpoint {
scanId: string; lastProcessedIndex: number; totalInvoices: number;
partialResults: any[]; sourceHash: string; timestamp: string;
}

export class FaultTolerantScanner {
private logger: winston.Logger;
private checkpointPath: string;
constructor(private readonly outputDir: string) {
this.checkpointPath = path.join(outputDir, '.checkpoint.json');
this.logger = winston.createLogger({
level: 'info',
format: winston.format.combine(
winston.format.timestamp(),
winston.format.json()
),
 transports: [
new winston.transports.File({
filename: path.join(outputDir, `healthledger-${Date.now()}.log`)
}),
],
});
this.setupGracefulShutdown();

```

```
}
// READ-ONLY: never write to source – only read
async loadSourceData(sourceFile: string): Promise<any[]> {
  const buffer = fs.readFileSync(sourceFile); // read-only access
  const hash = crypto.createHash('sha256').update(buffer).digest('hex');
  this.logger.info({ event: 'SOURCE_LOADED', sourceFile, hash });
  return { data: JSON.parse(buffer.toString()), hash };
}
// CHECKPOINT: save progress after each batch
saveCheckpoint(cp: ScanCheckpoint): void {
  const tmp = this.checkpointPath + '.tmp';
  fs.writeFileSync(tmp, JSON.stringify(cp, null, 2));
  fs.renameSync(tmp, this.checkpointPath); // atomic swap – no partial writes
  this.logger.info({ event: 'CHECKPOINT_SAVED', index: cp.lastProcessedIndex });
}
// RESUME: start from last checkpoint
loadCheckpoint(): ScanCheckpoint | null {
  if (!fs.existsSync(this.checkpointPath)) return null;
  return JSON.parse(fs.readFileSync(this.checkpointPath, 'utf-8'));
}
// INTEGRITY: verify source hash before and after scan
verifyIntegrity(sourceFile: string, expectedHash: string): boolean {
  const buffer = fs.readFileSync(sourceFile);
  const hash = crypto.createHash('sha256').update(buffer).digest('hex');
  const ok = hash === expectedHash;
  if (!ok) {
    this.logger.error({ event: 'INTEGRITY_VIOLATION', sourceFile, expected: expectedHash, got: hash });
    this.escalate('CRITICAL', 'Source data modified during scan – aborting');
  }
  return ok;
}
// Requirement 3 – Error Escalation Protocol
escalate(level: 'INFO'|'WARN'|'MEDIUM'|'CRITICAL', message: string, context?: any): void {
  this.logger.log(level === 'INFO' ? 'info' : level === 'WARN' ? 'warn' : 'error',
    { event: 'ESCALATION', level, message, context });
  if (level === 'MEDIUM') {
    // send notification to development team (email/API)
    this.sendNotification(message, level);
  }
  if (level === 'CRITICAL') {
    this.sendNotification(message, level, context);
    throw new Error(`CRITICAL: ${message}`);
  }
}
private sendNotification(message: string, level: string, context?: any): void {
  this.logger.info({ event: 'NOTIFICATION_SENT', level, message });
  // POST to monitoring webhook – placeholder for production implementation
}
// Requirement 2 – Remote Access: time-limited session key
generateRemoteSessionKey(durationMinutes: number = 60): string {
  const expiry = Date.now() + durationMinutes * 60 * 1000;
  const key = crypto.randomBytes(32).toString('hex');
  const keyData = { key, expiry, created: new Date().toISOString() };
  fs.writeFileSync(path.join(this.outputDir, 'remote-session.json'),
```

```
JSON.stringify(keyData, null, 2));
this.logger.info({ event: 'REMOTE_SESSION_CREATED', expiry: new Date(expiry).toISOString() });
return key;
}
validateRemoteSession(providedKey: string): boolean {
const keyFile = path.join(this.outputDir, '.remote-session.json');
if (!fs.existsSync(keyFile)) return false;
const { key, expiry } = JSON.parse(fs.readFileSync(keyFile, 'utf-8'));
return providedKey === key && Date.now() < expiry;
}
private setupGracefulShutdown(): void {
['SIGINT', 'SIGTERM'].forEach(signal => {
process.on(signal, () => {
this.logger.info({ event: 'GRACEFUL_SHUTDOWN', signal });
process.exit(0);
});
});
}
}
```

Code Listing 4 — lib/rag-compliance.ts (RAG Pipeline)

```
// lib/rag-compliance.ts — RAG Pipeline for Cooperation Agreement Compliance
import Anthropic from '@anthropic-ai/sdk';
// 6-type healthcare clause taxonomy
type ClauseType = 'reimbursement' | 'fee_schedule' | 'coverage_limit' |
'exclusion' | 'bundling_rule' | 'dispute_resolution';
interface AgreementChunk {
id: string; hospitalId: string; agreementRef: string;
clauseType: ClauseType; text: string;
embedding: number[]; // 1536-dim
}
const client = new Anthropic();
export async function embedText(text: string): Promise<number[]> {
// Generate 1536-dim embedding via Anthropic embeddings API
// Placeholder: actual implementation uses text-embedding-3-small or equivalent
return Array.from({ length: 1536 }, () => Math.random() * 2 - 1);
}
export async function storeAgreement(
hospitalId: string, agreementRef: string,
chunks: string[], clauseTypes: ClauseType[]
): Promise<void> {
for (let i = 0; i < chunks.length; i++) {
const embedding = await embedText(chunks[i]);
// Store in Pinecone/pgvector with metadata
console.log(`Stored chunk ${i}: ${clauseTypes[i]} — ${chunks[i].substring(0, 60)}...`);
}
}
export async function queryCompliance(
hospitalId: string, invoiceContext: string
): Promise<string> {
const queryEmbedding = await embedText(invoiceContext);
// Retrieve top-k relevant agreement chunks from vector DB
const relevantClauses = ['...retrieved clause 1...', '...retrieved clause 2...'];
// Inject into AI agent context for compliance analysis
const response = await client.messages.create({
model: 'claude-opus-4-8',
max_tokens: 2000,
messages: [{
role: 'user',
content: `Agreement clauses: ${relevantClauses.join('
')}
Invoice: ${invoiceContext}
Check compliance.`
}],
});
return (response.content[0] as any).text;
}
```

Code Listing 5 — Frontend (page.tsx, next.config.js, CNAME, globals.css)

```
// app/page.tsx — Next.js 14 Main Page (excerpt — brand content)
const PILLARS = [
'Cooperation Agreement Mapping',
'Accounting Compliance',
'Hospital-Specific AI',
```

```
'On-Site Methodology',
]
const LINES = [
  [{ text: 'The' }, { text: 'intelligence', cls: 'word-italic-sage' }],
  [{ text: 'healthcare' }, { text: 'accounting' }],
  [{ text: 'has been' }, { text: 'waiting for.', cls: 'word-italic-gold' }],
]
// → renders hero, ProblemHook, SectionHow, SectionContact, footer
// components/SectionHow.tsx — Three-Step Method
// Step 1: Agreement Ingestion — upload your cooperation agreements
// Step 2: Domain Training — AI trained on your specific rules
// Step 3: Live Compliance — real-time billing verification
// components/ProblemHook.tsx — Sequential problem display (1900ms intervals)
// "A hospital has 200+ cooperation agreements."
// "Each interpreted differently by accounting."
// "One clause. Thousands in exposure."
// CNAME → healthledgerai.com
// next.config.js → output: 'export', trailingSlash: true, images: {unoptimized: true}
// app/globals.css — CSS custom properties:
// --cream: #F4EFE4 --gold: #A08558 --sage: #4A6B5A
// --ivory: #EDE7D8 --blush: #D4A9A0 --ink: #1E1A14
// app/layout.tsx — Google Fonts: Cormorant Garamond + Jost
// metadata: { title: 'HealthLedger AI', description: 'Healthcare Accounting Intelligence' }
```

8. Academic Research Foundation

HealthLedger AI is grounded in a rigorous academic literature reviewed in the inventor's doctoral research proposal. The following evidence base documents: (a) the scope of the problem the invention addresses; (b) the absence of prior solutions; and (c) the academic validity of the detection methods employed. All sources below are cited with complete DOIs or verified URLs as they appear in the inventor's PhD Exposé and Scopus Literature Tracker.

8.1 Theoretical Frameworks Underpinning the Invention

Accounting Information Systems (AIS) Theory (Romney & Steinbart, 2021): Provides the foundational framework for understanding how AI transforms financial data collection, processing, and reporting — specifically data quality, process integrity, internal control effectiveness, and audit trail completeness. HealthLedger AI's alert lifecycle (OPEN → PENDING → RESOLVED → DISMISSED) and dated report archive directly implement these AIS constructs.

Human Error Theory & Cognitive Load (V■rzaru, 2024): AI adoption in accounting significantly reduces error frequency and improves fraud detection by automating cognitive tasks that cause human error — manual data entry, multitasking, time pressure. HealthLedger AI's 12 detection features target precisely these error sources.

Transaction Cost Economics (TCE) (Tenhunen, 2025): AI-driven accounting reduces transaction costs by automating repetitive tasks, minimising error-correction cycles, and streamlining compliance. The on-site discovery methodology and custom rules engine create a defensible cost structure (switching costs = competitive moat).

8.2 Empirical Evidence for the Problem

- **Billing error rates:** 7%–26% in manual hospital billing (Nimkar et al., 2024; V■rzaru, 2024); up to 80% of medical bills contain at least one error (Medical Billing Advocates of America, cited in Locumtenens.com, 2024)
- **AI error reduction:** More than 50% reduction in accounting errors with intelligent systems (Cunha, 2022; Wang, Zhang & Han, 2025)
- **Claim denials:** AI-enabled systems reduce claim denials by 15–30% and accelerate reimbursement cycles (mamun et al., 2025; Adelake & Ajayi, 2024)
- **Cost-effectiveness:** Most AI interventions in healthcare are cost-effective or cost-saving (Arab & Moosa, 2025 — Q1 journal, 95th percentile)
- **Fraud scale:** \$68 billion annually in US healthcare billing fraud (3–10% of total expenditure) (IJSRA, 2024)
- **Ghost patients:** \$61M Medicare fraud via ghost patients, Michigan, 2023 (Halunen Law, 2024)
- **RPA adoption:** Significant and accelerating uptake of automation in billing, claims, payment reconciliation (Nimkar et al., 2024)
- **AI fraud detection techniques:** Deep neural networks, XAI, federated learning increasingly effective for billing irregularities (Adelakun et al., 2024)

8.3 The Research Gap HealthLedger AI Addresses

A bibliometric analysis of 256 peer-reviewed articles (2015–2025) confirms AI as transformative in accounting and finance (Sanz Martín et al., 2025). An umbrella review of AI in hospital settings (Clark et al., 2025) explicitly identifies the absence of a comprehensive review synthesising all hospital financial management functions — confirming no prior system exists that treats hospital financial management as a unified domain.

Ma (2022) offers only a narrative treatment; Al-Hattami et al. (2025) confirm a critical research gap in AI adoption within healthcare finance. HealthLedger AI directly fills this gap.

The Cypriot and Greek contexts are particularly underresearched: the EU Digital Decade 2024 Country Report documents only 49.5% basic digital skills in Cyprus (below EU average 55.6%), with patient records largely paper-based. The inventor's PhD will provide the first systematic quantitative dataset on AI adoption in Cypriot hospital accounting — with HealthLedger AI as the operational prototype.

9. Prior Art Analysis

General Healthcare Billing Software

Examples: Epic, Meditech, Cerner

Key distinction: Handle billing workflow and EHR integration. Do NOT include: cooperation agreement mapping via RAG, on-site immersion methodology, 12-feature multi-type fraud detection, or fault-tolerant scan architecture.

Generic AI Fraud Detection

Examples: Palantir, DataRobot

Key distinction: General-purpose anomaly detection. Not hospital-specific, not cooperation-agreement-aware, no on-site discovery methodology, no academic PhD research grounding.

Revenue Cycle Management Tools

Examples: Waystar, Experian Health

Key distinction: Automate claims submission and denial management. No AI agent with persistent compliance memory, no RAG cooperation agreement mapping, no on-site immersion.

Accounting AI Tools

Examples: Sage Intacct AI, QuickBooks AI

Key distinction: General SME accounting. Not hospital-specific, no healthcare billing code logic, no cooperation agreement layer, no identity collision detection.

Academic Prototypes

Examples: Literature survey (Adelakun et al., 2024; Clark et al., 2025)

Key distinction: Documented but not commercially available: deep neural networks for billing fraud, XAI for audit — none with the combined cooperation agreement + on-site + fault-tolerant architecture of HealthLedger AI.

10. Draft Patent Claims

The following claims are preliminary drafts for attorney review. Five independent claims cover: (1) the 12-feature billing fraud detection system; (2) the fault tolerance architecture; (3) the AI compliance method; (4) the on-site immersion methodology; (5) the animated UI. Eleven dependent claims add specificity. All claims are based on the current production code and developer journal as of 28 March 2026.

Claim 1 (INDEPENDENT): Computer-implemented system for hospital billing fraud detection

A computer-implemented system for automated detection of billing fraud and accounting errors in hospital financial systems, comprising: (a) a 12-feature detection engine configured to identify duplicate invoices using vendor/amount/date fingerprinting with set-based $O(1)$ lookup; identity collisions between patients sharing name and date of birth with parental name secondary verification; upcoded procedures exceeding per-code maximum rates with overcharge percentage calculation; phantom billing through patient record and appointment cross-referencing; unbundled procedure codes through bundle-rule matching on grouped patient-date combinations; accumulated credit balances with 180-day regulatory threshold tracking; surprise and out-of-network billing including detection of doctors unregistered in any insurance network; and automated report generation in both .txt and .xlsx formats; (b) a risk-scoring module classifying alerts as LOW, MEDIUM, HIGH, or CRITICAL; (c) an alert lifecycle manager with states OPEN, PENDING, RESOLVED, and DISMISSED; (d) a CFO dashboard presenting totals, money protected, and money at risk; wherein the system processes real hospital invoice data from .xlsx files using pandas and returns results as HealthLedger_Alerts.xlsx.

Claim 2 (INDEPENDENT): Fault-tolerant hospital system scanning architecture

A fault-tolerant scanning architecture for use with live hospital information systems, comprising: (a) read-only access enforcement ensuring the source hospital database is never modified; (b) progressive checkpoint saving with atomic file swap to preserve partial results on crash; (c) checkpoint-and-resume capability starting subsequent scans from the last successful position; (d) transaction-safe write operations using atomic rename to prevent partial file writes; (e) SHA-256 hash verification of source data before and after each scan with automatic halt and escalation on integrity violation; (f) continuous structured logging via Winston with per-row processing records, timestamps, and error context; (g) a three-tier escalation protocol distinguishing silent minor errors, notification-triggering medium errors, and immediate-alert critical failures; and (h) a time-limited cryptographic remote session key for permission-based remote diagnostics and repair.

Claim 3 (INDEPENDENT): Method for AI-powered cooperation agreement compliance

A method for real-time hospital billing compliance using AI and cooperation agreement mapping, the method comprising: (a) chunking and embedding hospital-specific cooperation agreements as 1536-dimensional vectors with a 6-type healthcare clause taxonomy (reimbursement, fee_schedule, coverage_limit, exclusion, bundling_rule, dispute_resolution); (b) storing embeddings in a vector database; (c) at scan time, embedding the invoice context and retrieving the top-k most relevant agreement clauses via cosine similarity; (d) injecting retrieved clauses into a persistent AI agent (claude-opus-4-8 via Anthropic Managed Agents API) for compliance analysis with streaming output; (e) returning a compliance determination with supporting clause citations; wherein the agent's AGENT_ID is stored durably for reuse across sessions.

Claim 4 (INDEPENDENT): On-site immersion methodology for custom rule discovery

A methodology for hospital accounting intelligence deployment comprising: (a) a 2–3 day on-site discovery session with the hospital accounting department prior to system deployment; (b) systematic documentation of the hospital's cooperation agreements, insurer contracts, surgeon fee schedules, and institutional billing patterns; (c) identification of hidden compliance patterns not visible in written agreements; (d) encoding of discovered rules into a configurable Custom Rules Engine specific to that hospital; wherein the resulting rule set creates a structural switching cost that is unique to each institution and not reproducible without repeating the on-site process.

Claim 5 (INDEPENDENT): Diamond glitter particle burst UI animation

A user interface comprising a four-phase animated entry sequence: Phase 1 (IDLE, 0–2600ms) — a dense silver/cool-palette diamond glitter particle sphere breathing on a dark (#15120E) backdrop with unit-sphere positioning and sinusoidal opacity twinkle; Phase 2 (BURST, 2600–5400ms) — particles launched radially outward using unit-sphere direction vectors with momentum accumulation; Phase 3 (HOLD, 5400–6400ms) — scattered glitter constellation lingering on dark backdrop; Phase 4 (FADE, 6400–7800ms) — dark overlay dissolves to reveal the branded website; wherein the animation plays automatically on page load and requires no user interaction to advance.

Claim 6 (dep. on 1): The system of Claim 1, further comprising a parental name verification prompt that is automatically triggered when two patients share the same first name, last name, and date of birth, blocking invoice processing until the secondary identifier is confirmed.

Claim 7 (dep. on 1): The system of Claim 1, wherein the Excel output file HealthLedger_Alerts.xlsx contains columns for Invoice ID, Vendor, Amount, Date, Alert Type, and Overcharge Percentage, operable by any standard spreadsheet application without additional software.

Claim 8 (dep. on 1): The system of Claim 1, wherein the credit balance module calculates the number of days elapsed since invoice date and triggers a HIGH alert when the credit balance age approaches 180 days, the threshold for unclaimed property reporting obligations.

Claim 9 (dep. on 2): The fault-tolerant architecture of Claim 2, wherein the remote session key contains a cryptographic expiry timestamp and the validation function rejects the key after expiry, ensuring no permanent remote access capability persists after the repair session.

Claim 10 (dep. on 2): The fault-tolerant architecture of Claim 2, wherein the structured log file is self-contained and transmissible without technical expertise, enabling hospital staff to send the log file to the development team as a complete diagnostic record.

Claim 11 (dep. on 3): The method of Claim 3, wherein the vector database stores agreement chunks with hospitalId and agreementRef metadata, enabling per-hospital compliance isolation and preventing cross-hospital data leakage.

Claim 12 (dep. on 3): The method of Claim 3, wherein the 6-type clause taxonomy classifies each agreement chunk at ingestion time, and retrieval can be filtered by clause type to retrieve only reimbursement clauses, only exclusion clauses, or only bundling rules as required.

Claim 13 (dep. on 4): The methodology of Claim 4, wherein the on-site discovery session is documented in a standardised format that serves simultaneously as the hospital's Custom Rules Engine configuration and as primary research data for the inventor's doctoral dissertation on AI adoption in hospital accounting.

Claim 14 (dep. on 5): The animation of Claim 5, wherein individual particles use unit-sphere coordinates (bx, by, bz) for position during the idle phase and direction during the burst phase, enabling smooth visual continuity between the breathing sphere and the radial burst without discontinuous position jumps.

Claim 15 (dep. on 1, 2, 3): The combined system of Claims 1, 2, and 3, further comprising a monitoring interface displaying which hospital installations are running normally, which have active warnings, and which are offline, updated via the fault tolerance logging infrastructure of Claim 2.

Claim 16 (dep. on 1, 4): The system of Claim 1, deployed according to the methodology of Claim 4, wherein the Custom Rules Engine contains hospital-specific rules discovered during the on-site session that are not available in any public or standardised billing code database and constitute a proprietary configuration unique to that institution.

11. Complete References

All references below are drawn from the inventor's PhD Exposé (March 2026) and Developer Journal (March 2026). Sources are given with complete DOIs or verified URLs. Journal quality indicators are from the Scopus Literature Tracker.

11.1 Academic References (Peer-Reviewed — with DOIs)

- Adelake, O., & Ajayi, S. (2024).** Transforming the Healthcare Revenue Cycle with Artificial Intelligence in the USA. *International Journal of Multidisciplinary Research and Growth Evaluation*, 5(3), 1069–1083. DOI/URL: <https://doi.org/10.54660/IJMRGE.2024.5.3.1069-1083>
- Adelakun, B., Onwubuariri, E., Adeniran, G., & Ntiakoh, A. (2024).** Enhancing fraud detection in accounting through AI: Techniques and case studies. *Finance & Accounting Research Journal*, 6(6), 978–999. DOI/URL: <https://doi.org/10.51594/farj.v6i6.1232>
- Al-Hattami, H. M., et al. (2025).** Bibliometric landscape of emerging technology in the accounting information systems field. *Cogent Business & Management*. [Scopus Q2, 70th percentile] DOI/URL: <https://doi.org/10.1080/23311975.2025.2573191>
- Arab, R. A., & Moosa, O. A. (2025).** Systematic review of cost-effectiveness and budget impact of artificial intelligence in healthcare. *npj Digital Medicine*, 8(1). [Scopus Q1, 95th percentile — elite journal] DOI/URL: <https://doi.org/10.1038/s41746-025-01722-y>
- Clark, S., et al. (2025).** An umbrella review of the facilitators and barriers to implementing AI solutions within hospital settings: through the lens of the NASSS framework. *openRxiv* (preprint). DOI/URL: <https://doi.org/10.1101/2025.06.19.25329916>
- Cunha, L. (2022).** Reducing Accounting Errors Through Intelligent Automation: A Literature Review and Practical Insights. *International Seven Journal of Multidisciplinary*, 1(2). DOI/URL: <https://doi.org/10.56238/isevmjv1n2-015>
- IJSRA — International Journal of Science and Research Archive. (2024).** Fraud detection in healthcare billing and claims. DOI/URL: <https://doi.org/10.30574/ijsra.2024.13.2.2606>
- Ma, M. (2022).** Research on the Development of Hospital Intelligent Finance Based on Artificial Intelligence. *Computational Intelligence and Neuroscience*, 2022, 1–6. DOI/URL: <https://doi.org/10.1155/2022/6549766>
- mamun, A., et al. (2025).** Optimizing Revenue Cycle Management in Healthcare: AI and IT Solutions for Business Process Automation. *The American Journal of Engineering and Technology*, 07(03), 141–162. DOI/URL: <https://doi.org/10.37547/tajet/Volume07Issue03-14>
- Nimkar, S., et al. (2024).** Increasing Trends of Artificial Intelligence With Robotic Process Automation in Health Care: A Narrative Review. *Cureus*, 16(9): e69680. [Scopus Q2, 72nd percentile] DOI/URL: <https://doi.org/10.7759/cureus.69680>
- Romney, M. B., & Steinbart, P. J. (2021).** *Accounting Information Systems* (15th ed.). Pearson. [Primary AIS theory anchor — data quality, process integrity, internal controls]
- Sanz Martín, L., Parra-Domínguez, J., Corchado, J. M., & Zafra-Gómez, E. (2025).** Recent evolution and growth of AI and advanced technologies in accounting and finance: systematic review and bibliometric analysis. *Spanish Journal of Finance and Accounting*. [Scopus Q2, 60th percentile] DOI/URL: <https://doi.org/10.1080/02102412.2025.2582120>
- Tenhunen, M.-L. (2025).** AI-Driven Management Accounting: A New Frontier in Strategic Finance. *Economics*, 14(4), 87–95. [TCE framework — proactive strategic finance] DOI/URL: <https://doi.org/10.11648/j.eco.20251404.11>
- V■rzaru, A. A. (2024).** Assessing the Transformative Impact of AI Adoption on Efficiency, Fraud Detection, and Skill Dynamics in Accounting Practices. *Journal of Risk and Financial Management*, 17(12), 577. [Scopus Q2, 65th percentile] DOI/URL: <https://doi.org/10.3390/jrfm17120577>

Wang, M., Zhang, X., & Han, X. (2025). AI Driven Systems for Improving Accounting Accuracy Fraud Detection and Financial Transparency. *Frontiers in Artificial Intelligence Research*, 2(3), 403–421. DOI/URL: <https://doi.org/10.71465/fair398>

Zhou, X., et al. (2025). The Impact of Artificial Intelligence on Accounting Information and Earnings Management: Bibliometric Analysis. *Journal of Risk and Financial Management*, 19(1), 90. DOI/URL: <https://www.mdpi.com/1911-8074/19/1/90>

11.2 Government & Institutional Reports

ASTP — Office of the Assistant Secretary for Technology Policy. (2024). Hospital Trends in the Use, Evaluation, and Governance of Predictive AI, 2023–2024. ASTP Health IT Data Brief. URL: <https://www.ncbi.nlm.nih.gov/books/NBK618497/>

European Commission. (2024). Cyprus 2024 Digital Decade Country Report. URL: <https://digital-strategy.ec.europa.eu/en/factpages/cyprus-2024-digital-decade-country-report>

WHO European Observatory on Health Systems and Policies. (2024). Cyprus: Health System Summary 2024. URL: <https://eurohealthobservatory.who.int/publications/i/cyprus-health-system-summary-2024>

11.3 Industry & Legal Sources (from Developer Journal)

Becker's Hospital Review. (2024). 5 Common Accounting Blunders Hospitals Can Avoid. URL: <https://www.beckershospitalreview.com/finance/5-common-accounting-blunders-hospitals-can-avoid.html>

Halunen Law. (2024). Upcoding and Unbundling Are Common Types of Healthcare Fraud. URL: <https://www.halunenlaw.com/upcoding-and-unbundling-types-of-healthcare-fraud>

Healthcare Innovation Group. (2024). Medical Billing Errors Are Alarming Common — and Patients Are Paying the Price. URL: <https://www.hcinnovationgroup.com/finance-revenue-cycle/article/21080693>

Locumtenens.com. (2024). 8 Common Mistakes in Healthcare Billing. URL: <https://www.locumtenens.com/news-and-insights/blog/8-common-mistakes-in-healthcare-billing>

Medical Billing Advocates of America. (2024). Estimate: up to 80% of medical bills contain at least one billing error. Cited in: Locumtenens.com (2024).

12. Abstract

HealthLedger AI is an AI-powered hospital billing fraud detection, accounting error identification, and cooperation agreement compliance system invented by Maria Polychroniadou, Ph.D. candidate at the European University Cyprus (EUC). Conceived on 14 March 2026 and reduced to practice on 17 March 2026 (Version 1.0), the system combines: a 12-feature TypeScript detection engine covering duplicate invoices, patient identity collisions, upcoding, phantom billing, unbundling, credit balance accumulation, and surprise/out-of-network billing; a persistent AI agent (claude-opus-4-8, Anthropic Managed Agents API) for compliance analysis; a RAG pipeline with Pinecone/pgvector vector database and 6-type healthcare clause taxonomy for cooperation agreement mapping; and a fault-tolerant Phase 2 architecture (designed 28 March 2026) with read-only access, checkpoint/resume, SHA-256 integrity verification, Winston logging, three-tier escalation, and time-limited cryptographic remote session keys. The system is deployed as a Next.js 14 static export at healthledgerai.com. The invention addresses a documented research gap (Clark et al., 2025; Al-Hattami et al., 2025): no prior system treats hospital financial management as a unified domain with cooperation-agreement-aware compliance, on-site discovery methodology, and production-grade fault tolerance. The research foundation is the inventor's PhD Exposé supervised by Prof. Alexios Kythreotis, EUC, with study sites at the American Medical Center (Nicosia, Cyprus) and the European Interbalkan Medical Center (Thessaloniki, Greece).

13. Inventor Declaration

I, Maria Polychroniadou, hereby declare that I am the sole inventor of the subject matter described in this disclosure. I first conceived of HealthLedger AI on Saturday, 14 March 2026, and first reduced it to practice on Tuesday, 17 March 2026 (Version 1.0 — 11 core detection features). The Excel integration (Feature 12) was completed on 17–18 March 2026. The Phase 2 fault tolerance and remote access architecture was designed and documented on 28 March 2026. This disclosure has been prepared in good faith and represents my best understanding of the novel subject matter as of the date of this document.

Inventor Name:	Maria Polychroniadou
Institution:	Ph.D. Candidate (proposed), Accounting & Finance, European University
Supervisor:	Cyprus Prof. Alexios Kythreotis, EUC
Email:	polychroniadou@gmx.de · info@healthledgerai.com
Date of Conception:	14 March 2026 (Saturday)
Date of Disclosure:	26 June 2026
Signature:	

Witnessed By:

Witness Date:

Notary / Stamp:

CONFIDENTIAL — ATTORNEY-CLIENT PRIVILEGED — FINAL v5 © 2026 HealthLedger AI · Maria Polychroniadou · All rights reserved.
healthledgerai.com · info@healthledgerai.com · Conceived: 14 March 2026 · Disclosed: 26 June 2026